HPCS
High Performance Computer System Lab

Flexon: A Flexible Digital Neuron for Efficient Spiking Neural Network Simulations

Dayeol Lee[†], Gwangmu Lee[*], Dongup Kwon[*],
Sunghwa Lee[*], Youngsok Kim[*], and Jangwoo Kim[*]

[*]Dept. of Electrical and Computer Engineering, Seoul National University
[†]Dept. of Electrical Engineering and Computer Sciences, University of California, Berkeley

Degeneration

Parkinson's Disease,
CJD, Dementia

Recognition

Object detection,
Classification

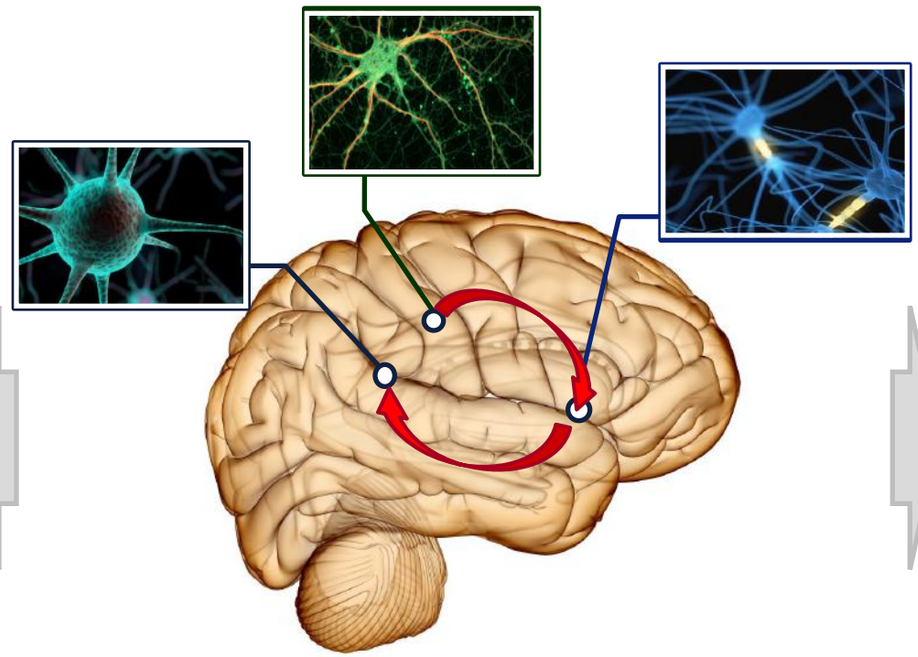

Consciousness

Morality,
Social Value

Emotion

Sympathy,
Happiness, Apathy

Neuroscience

The Study of **Neurons** and **Brains**

DNN did this

Degeneration

Parkinson's Disease,
CJD, Dementia

Recognition

Object detection,
Classification

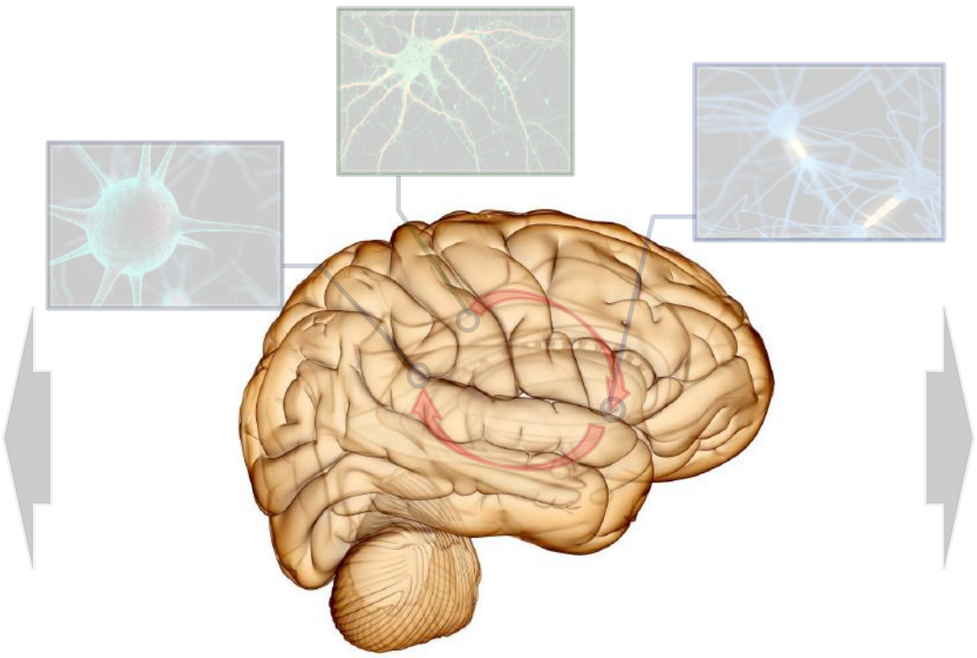

Consciousness

Morality,
Social Value

Emotion

Sympathy,
Happiness, Apathy

Neuroscience

The Study of **Neurons** and **Brains**

Unexplored yet

Modeling a Brain: Spiking Neural Network

How to compute spiking neural network?

```
for each time-step:
```

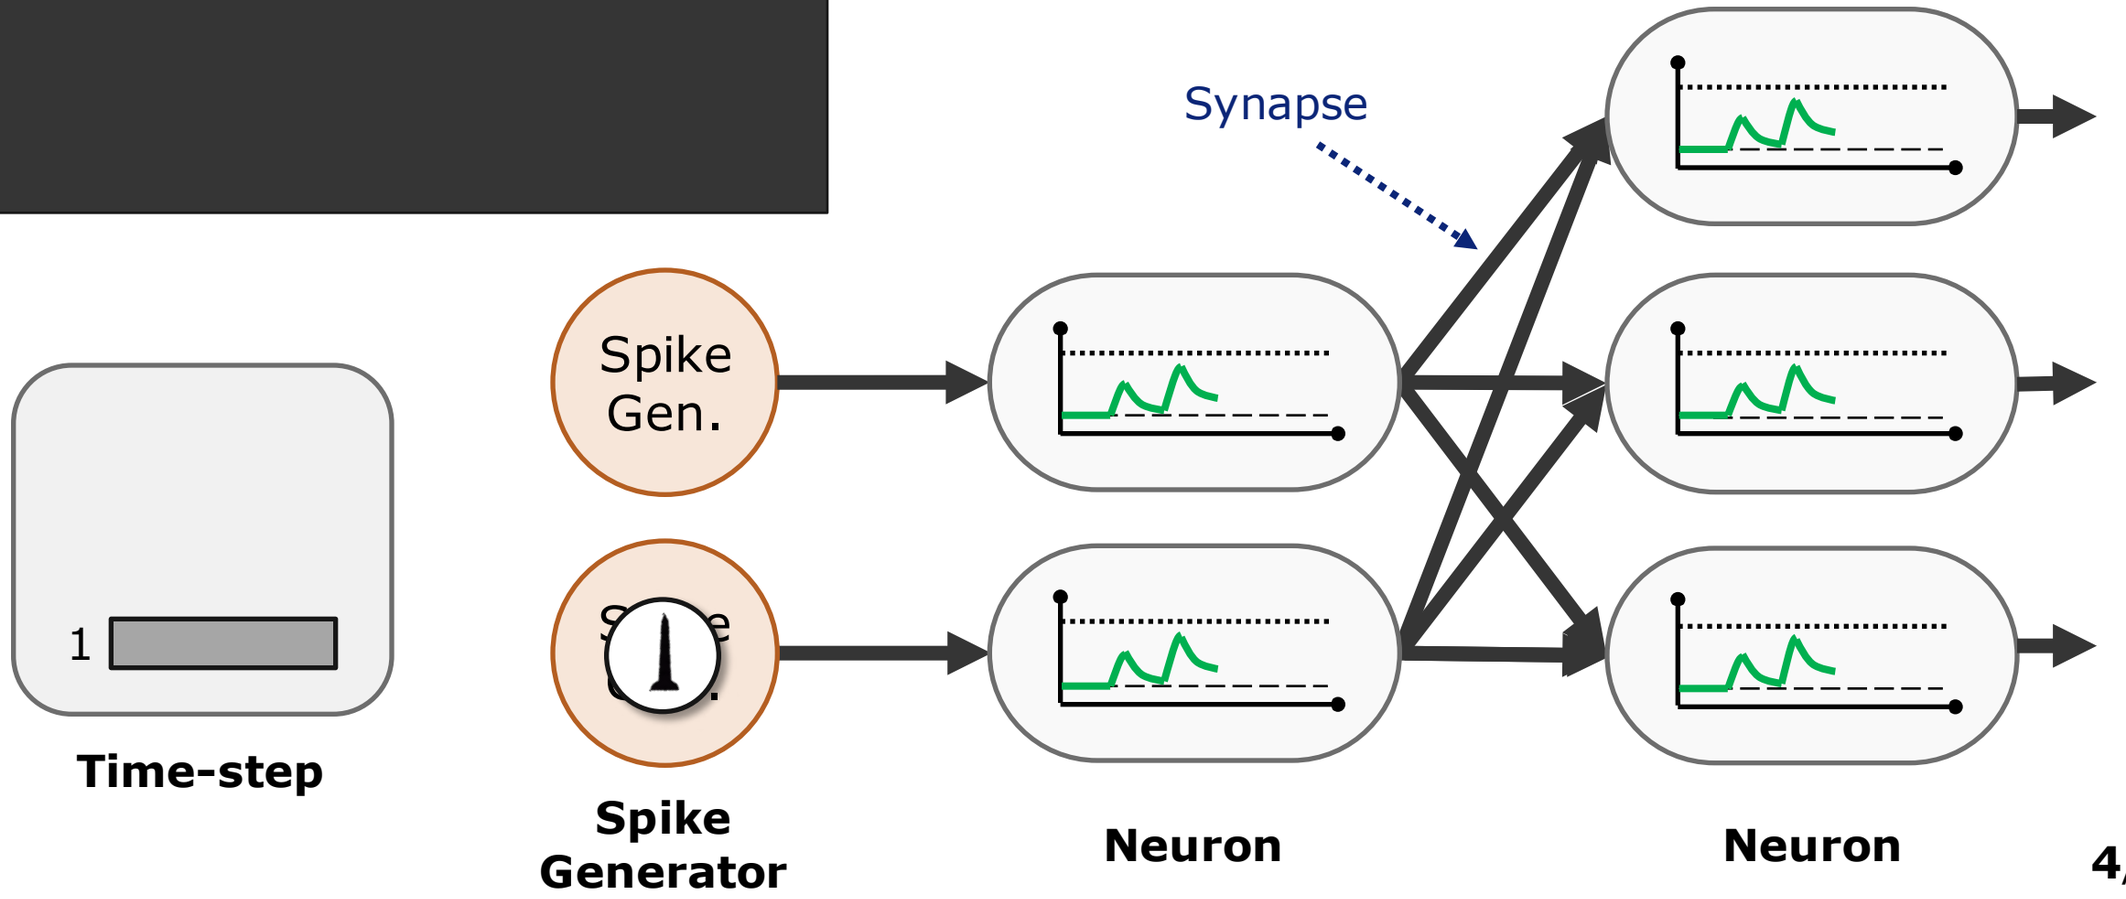

Modeling a Brain: Spiking Neural Network

How to compute spiking neural network?

```
for each time-step:
  1) stimulus generation
```

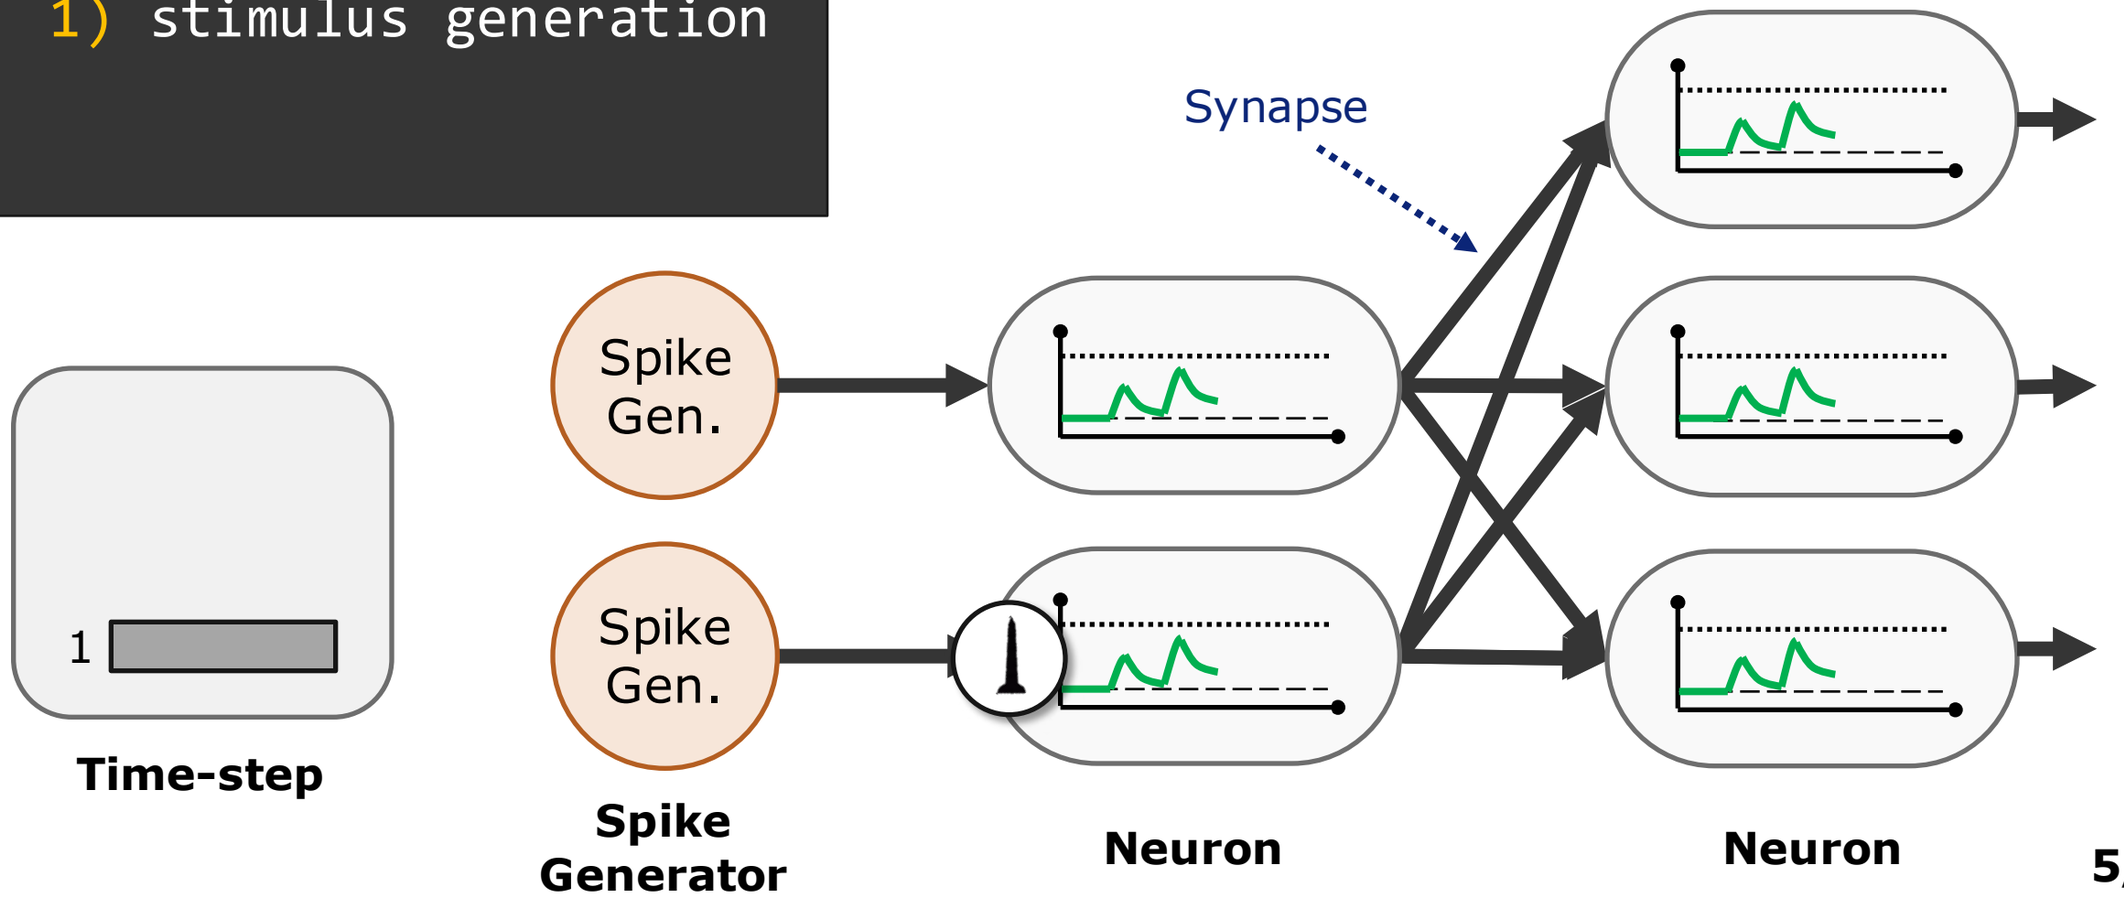

Modeling a Brain: Spiking Neural Network

How to compute spiking neural network?

```
for each time-step:
  1) stimulus generation
  2) neuron computation
```

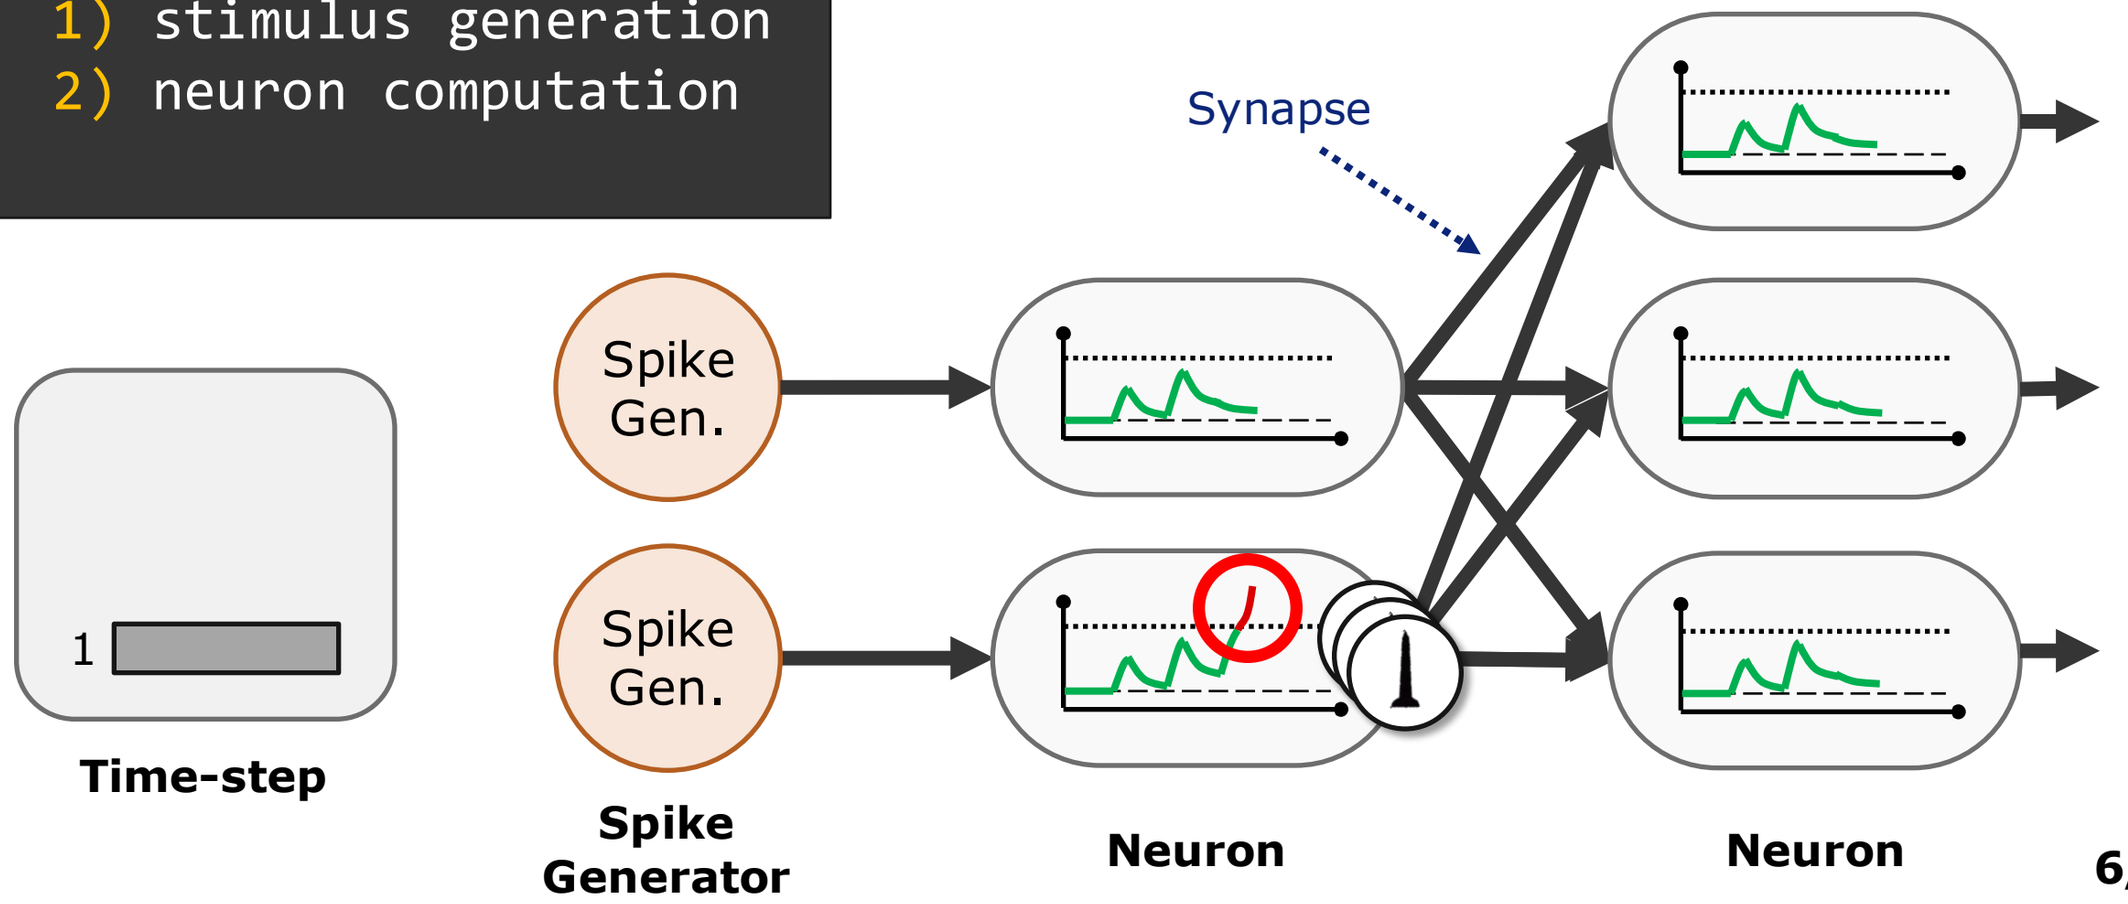

Modeling a Brain: Spiking Neural Network

How to compute spiking neural network?

```
for each time-step:
  1) stimulus generation
  2) neuron computation
  3) synapse calculation
```

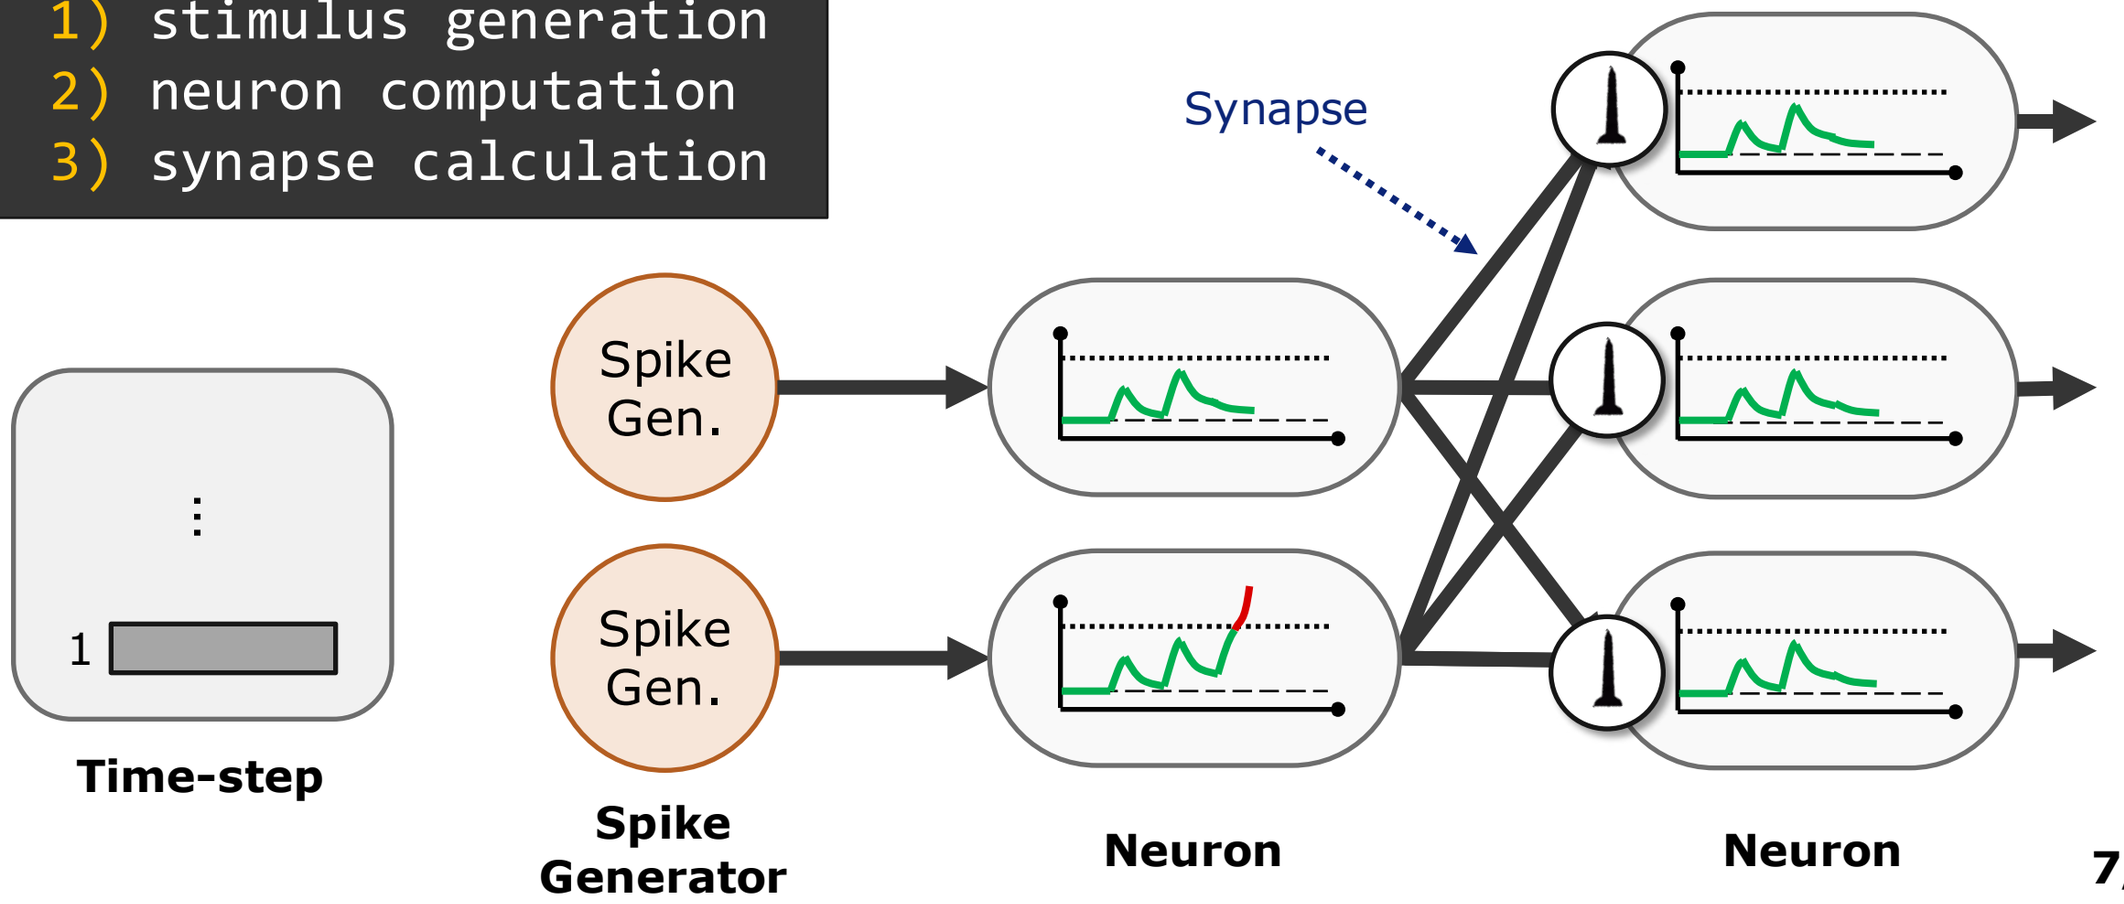

Where Does Time Go?

10 Representative Benchmarks on CPU/GPU

CPU: Intel Xeon E5-2630 v4 CPU (12-core, 2.2 GHz) / GPU: NVIDIA Titan X (Pascal) GPU

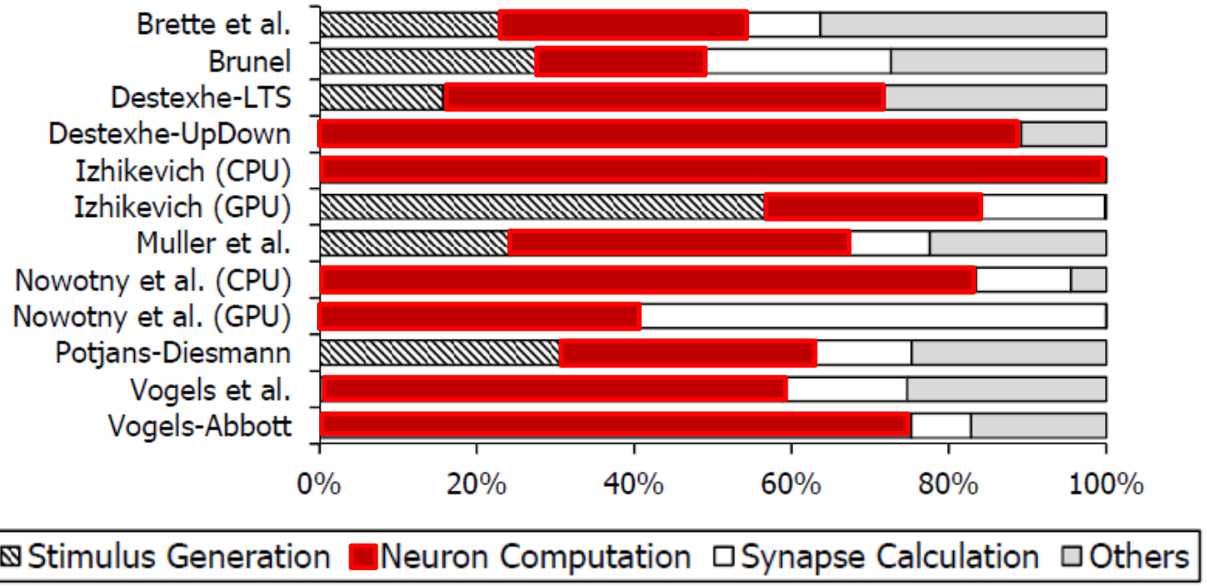

~50% of overheads coming from

Neuron Computation

How Does a Neuron Behave?

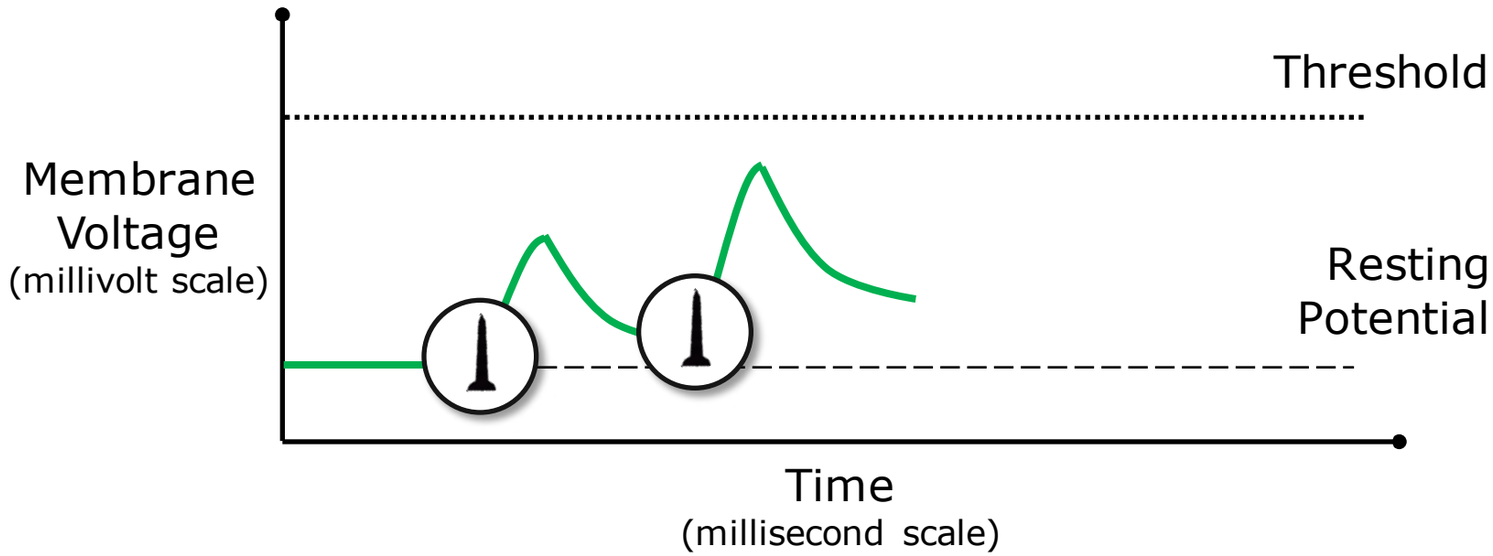

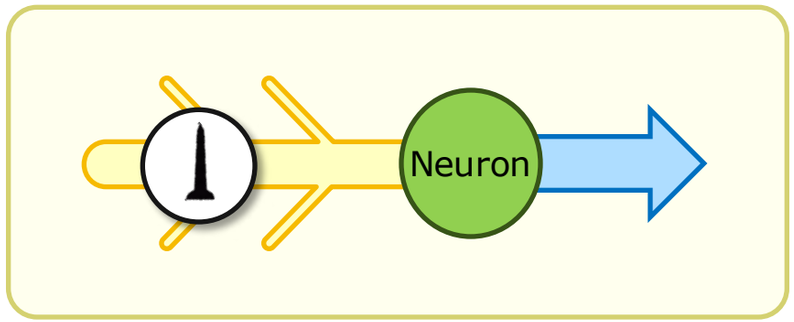

Biological Neuron

How Does a Neuron Behave?

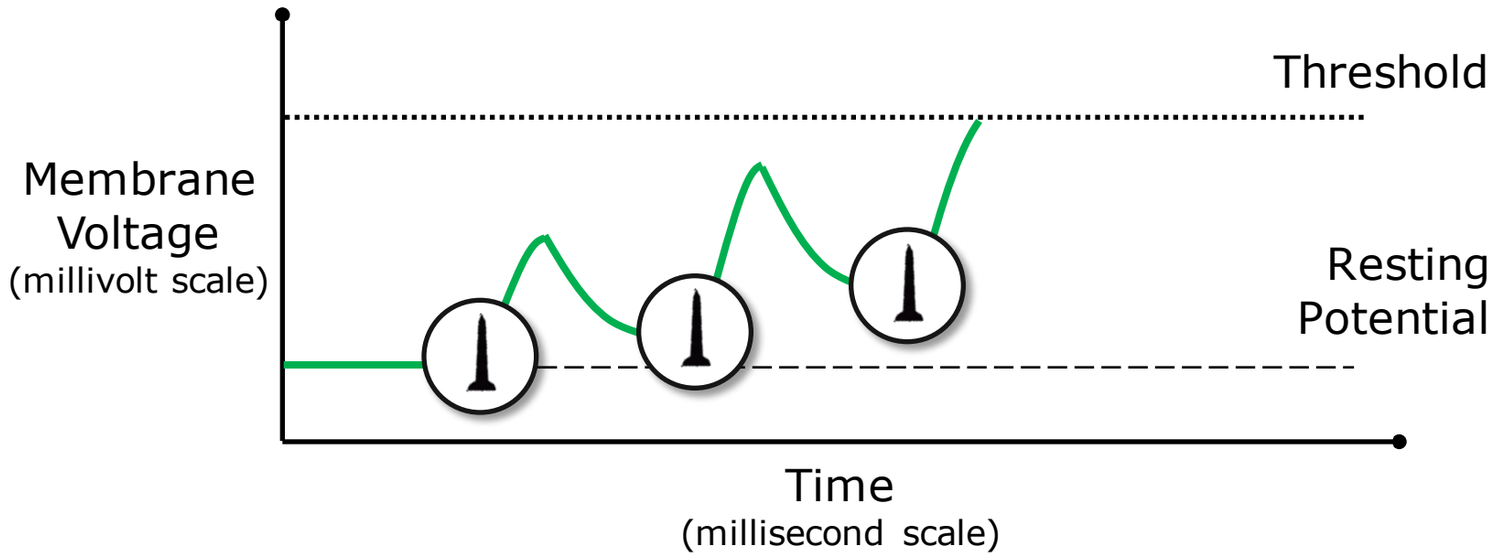

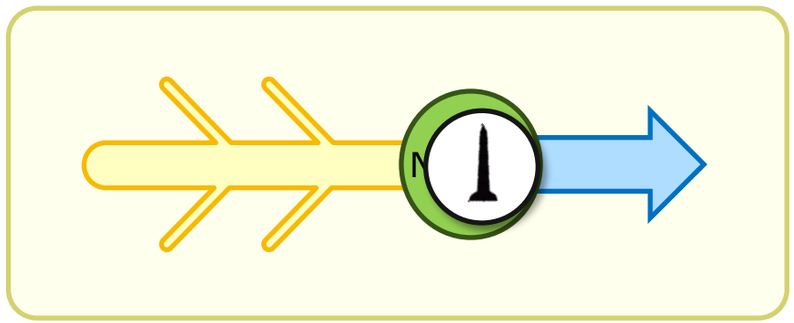

Biological Neuron

How Does a Neuron Behave?

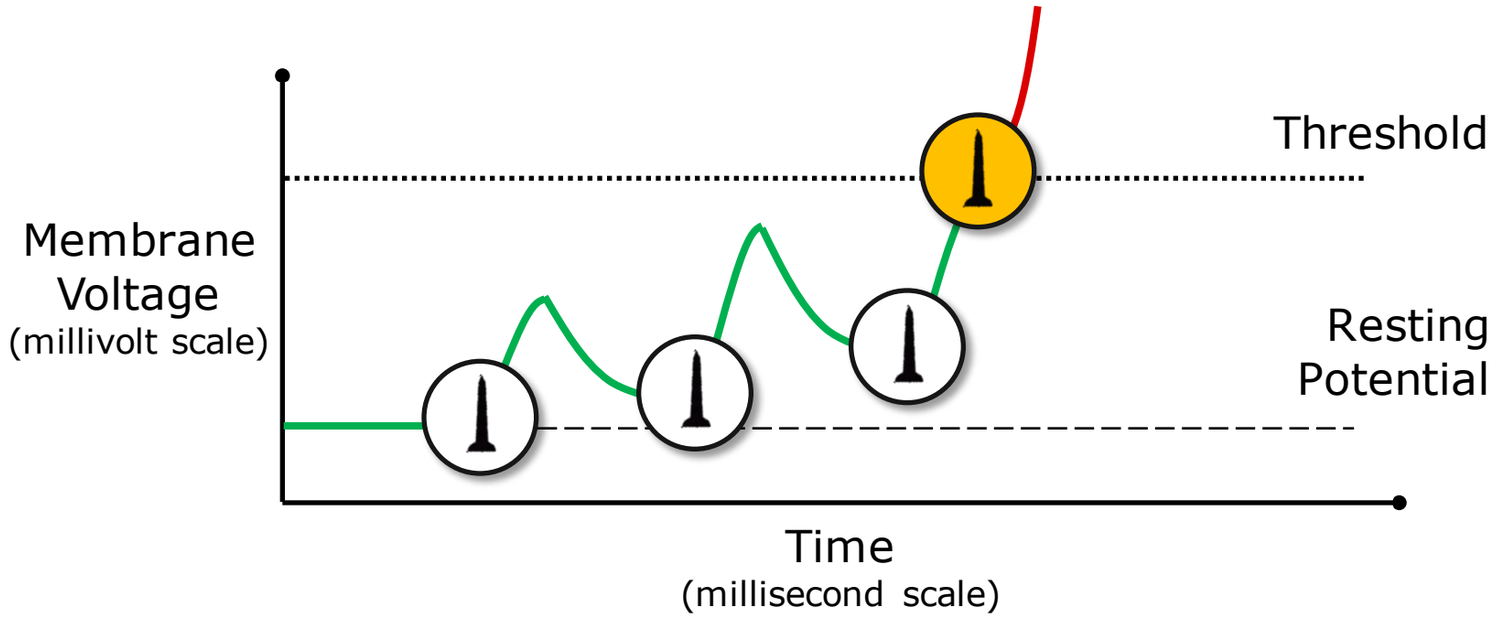

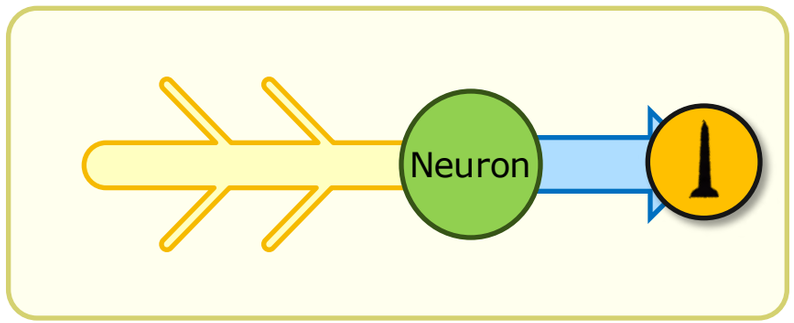

Biological Neuron

How Does a Neuron Behave?

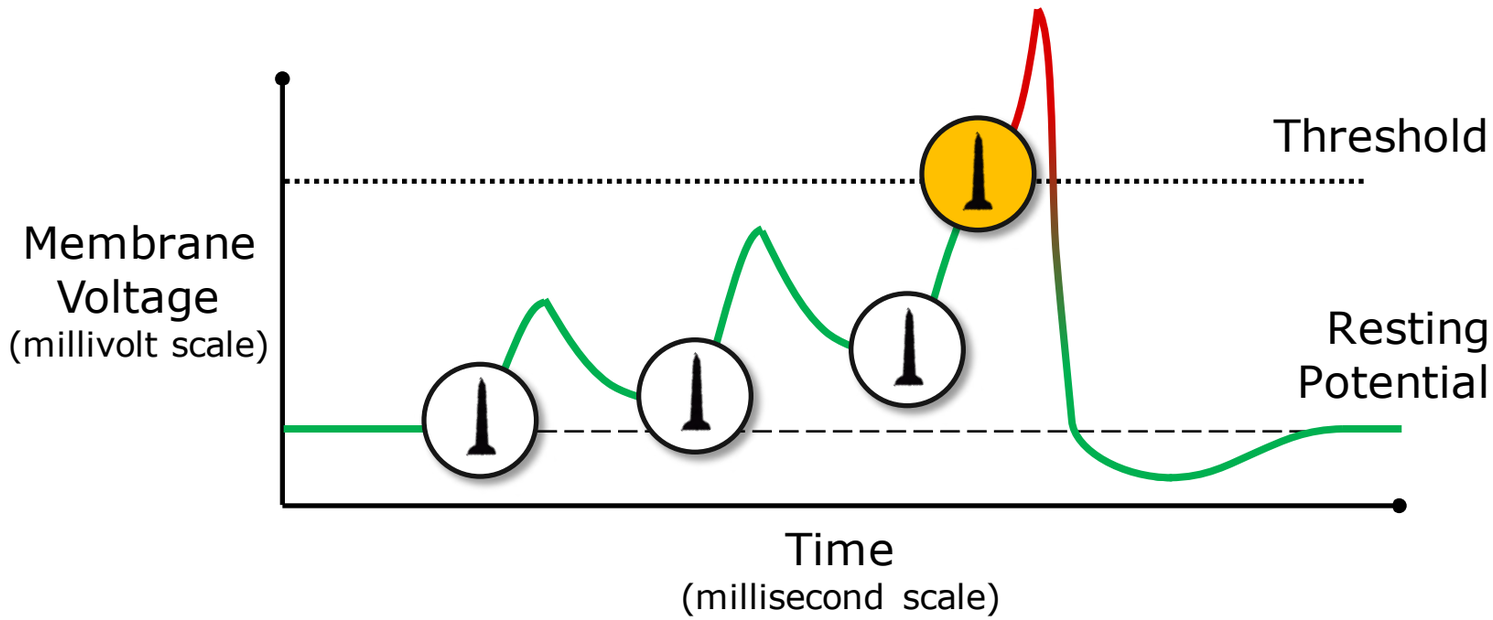

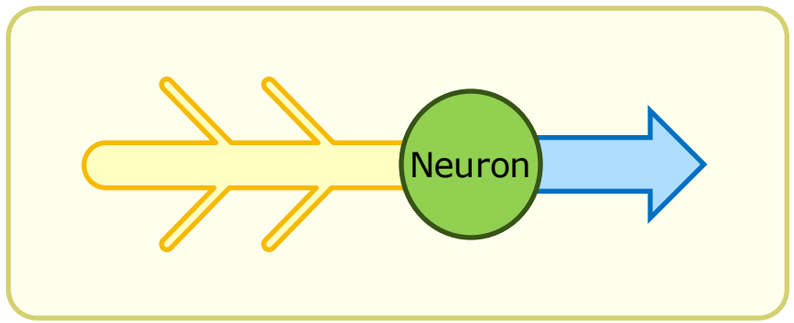

Biological Neuron

Various Neuron Behaviors

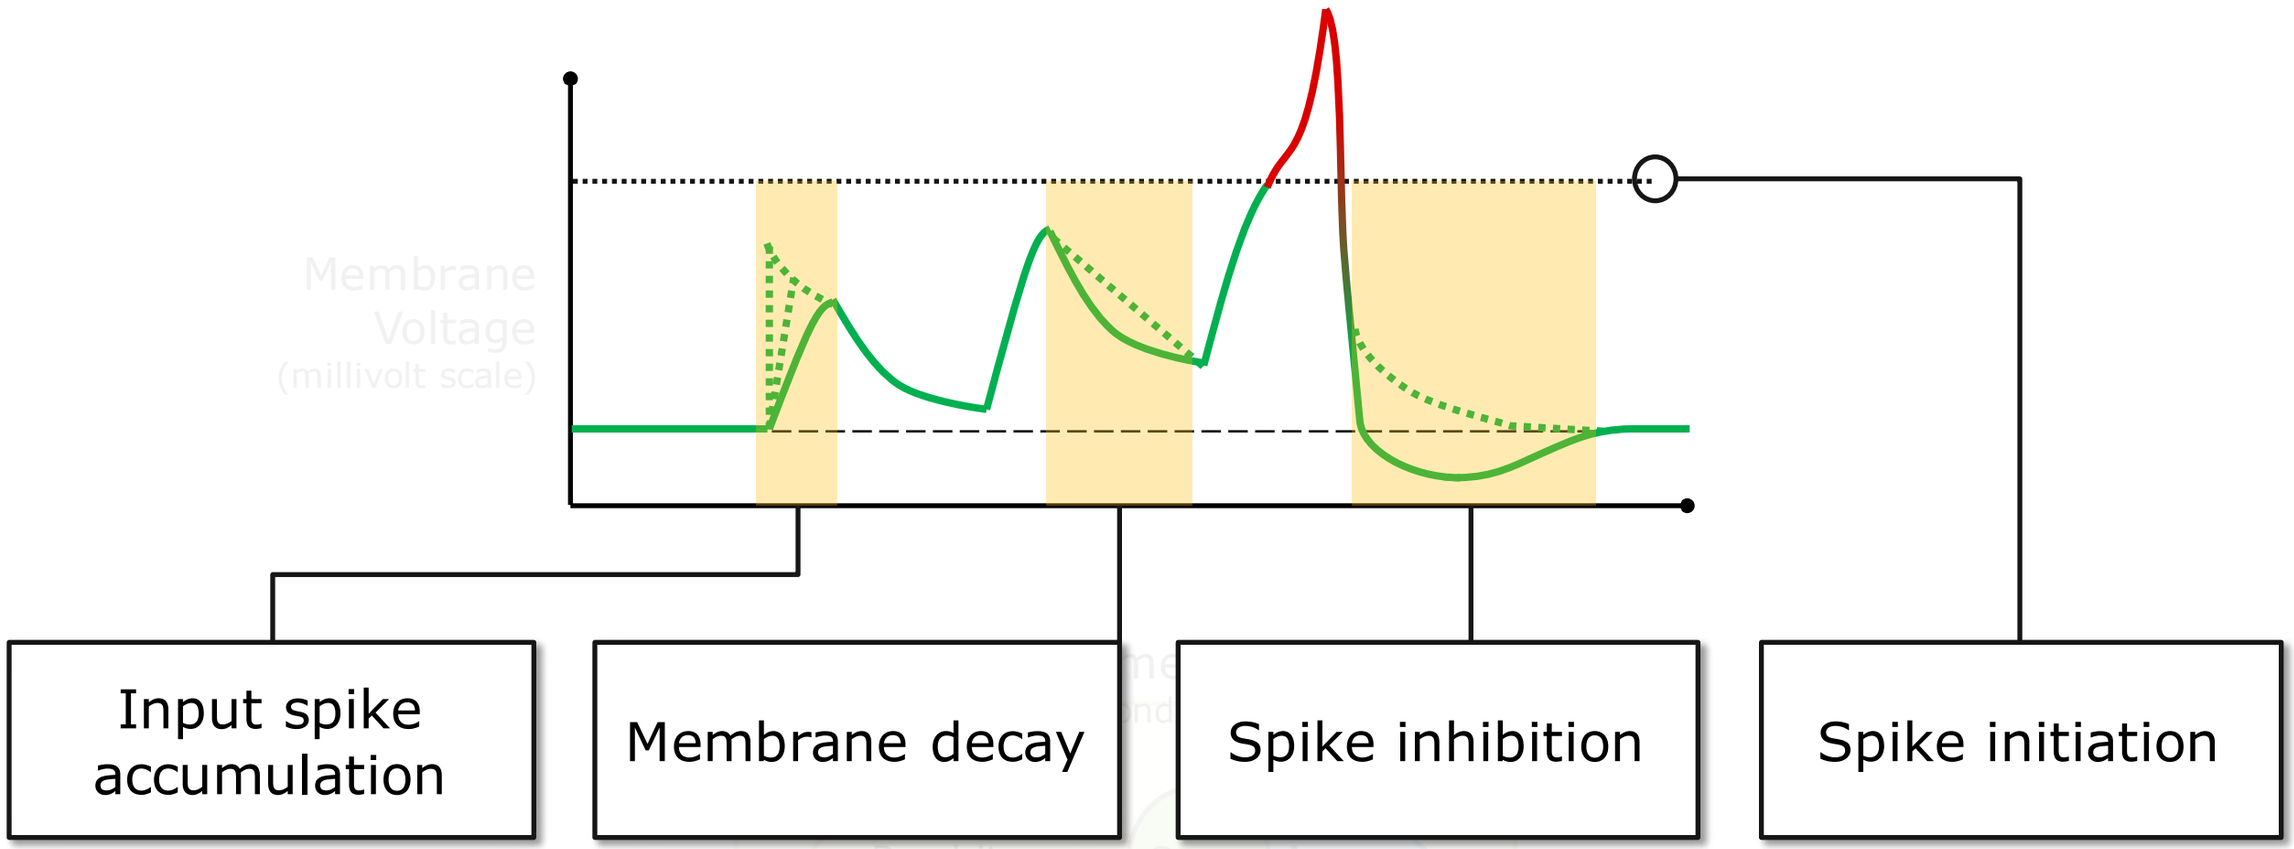

Tons of variants exist, depending on their **feature set.**

We need to support **various features** for **accurate** brain simulations.

Solutions and Limitations

Software Simulation		Custom Hardware Framework
✓	Flexibility	✗
✓	Accuracy	✗
✗	High Performance	✓
✗	Low Energy	✓

Design Goals & Key Ideas

Goal 1	High Performance	→ Hardware-based
Goal 2	High Flexibility	→ Feature-driven
Goal 3	Low cost	→ Spatially-folded

Neuron Feature #1: **Input Spike Accumulation**

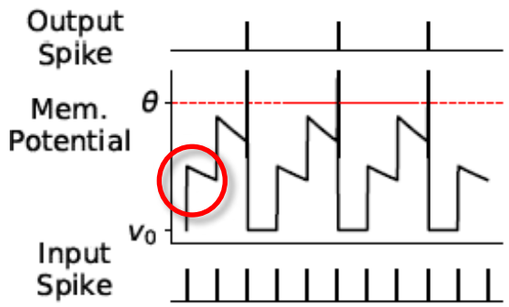

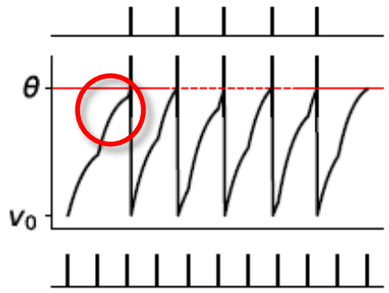

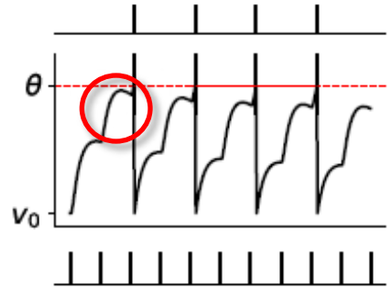

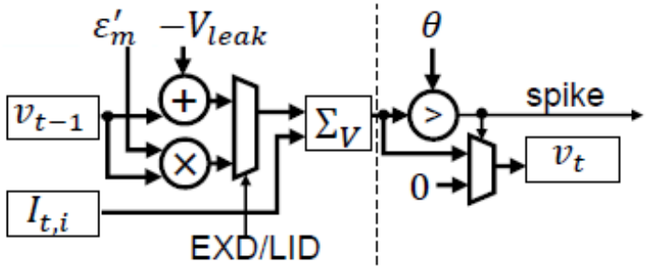

Current-based

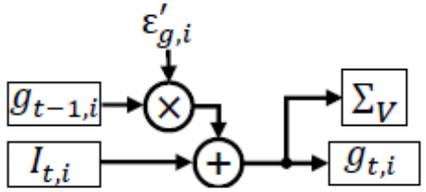

Conductance-based
(Exponential-shaped)

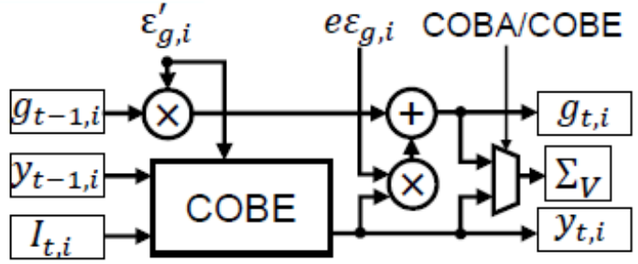

Conductance-based
(Alpha function-shaped)

Neuron Feature #1: **Input Spike Accumulation**

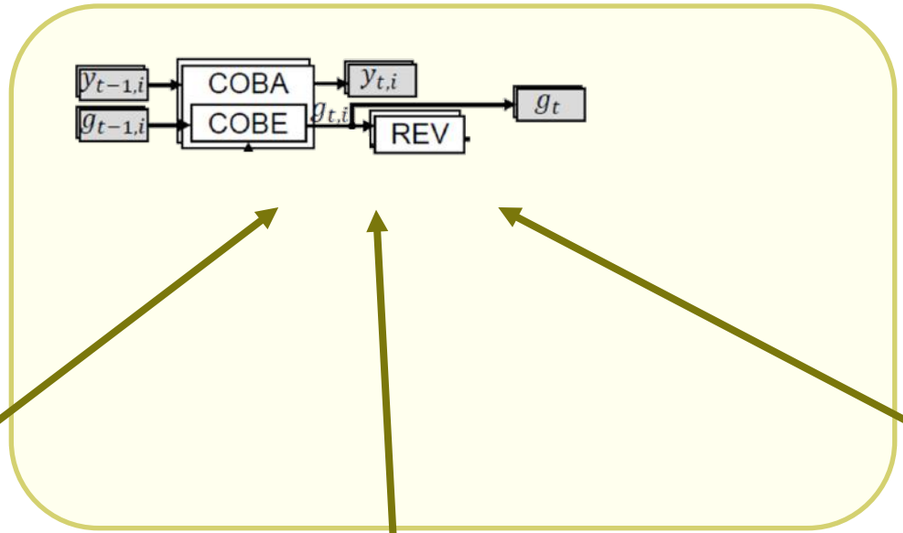

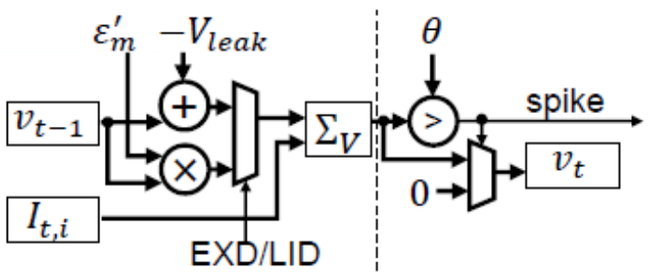

Current-based

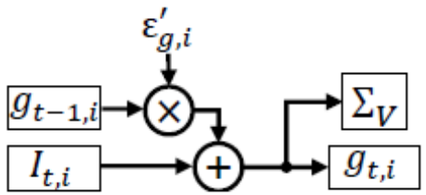

Conductance-based
(Exponential-shaped)

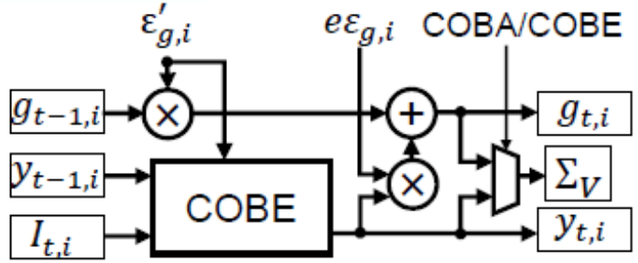

Conductance-based
(Alpha function-shaped)

Neuron Feature #2: **Spike Initiation**

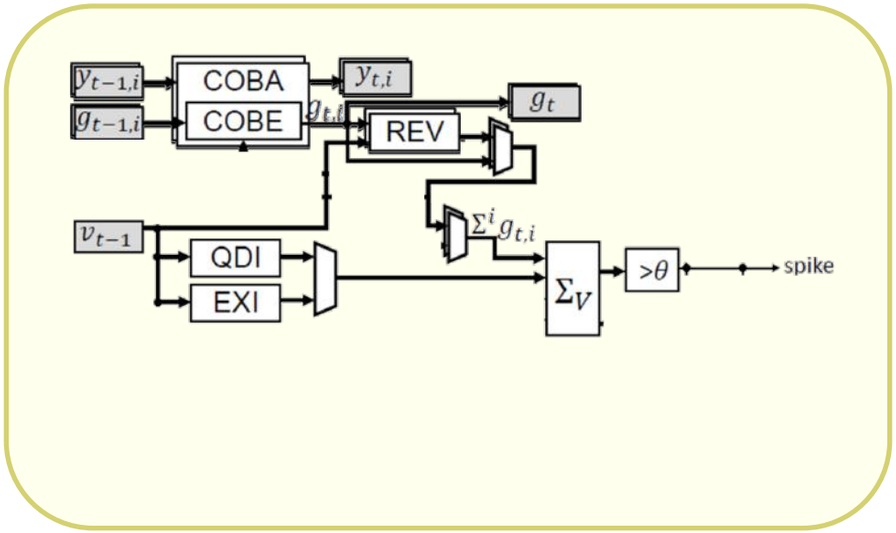

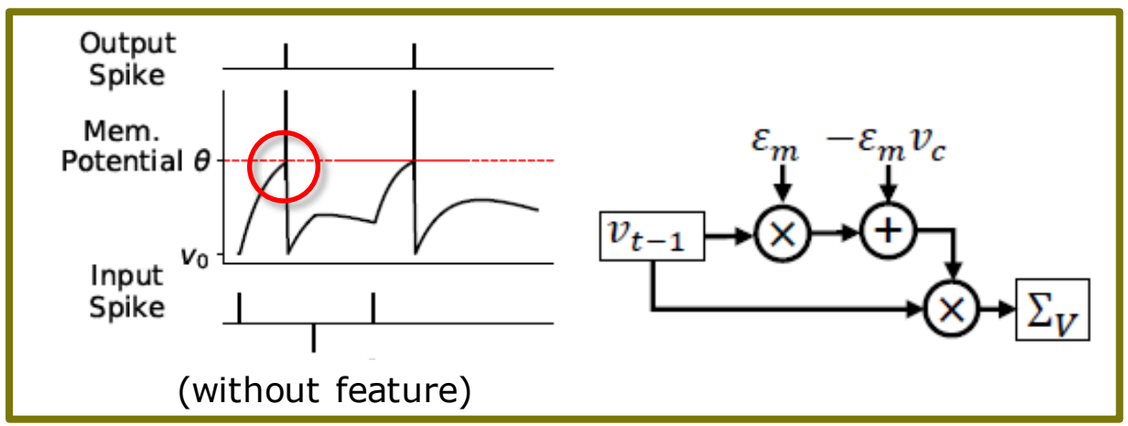

Quadratic

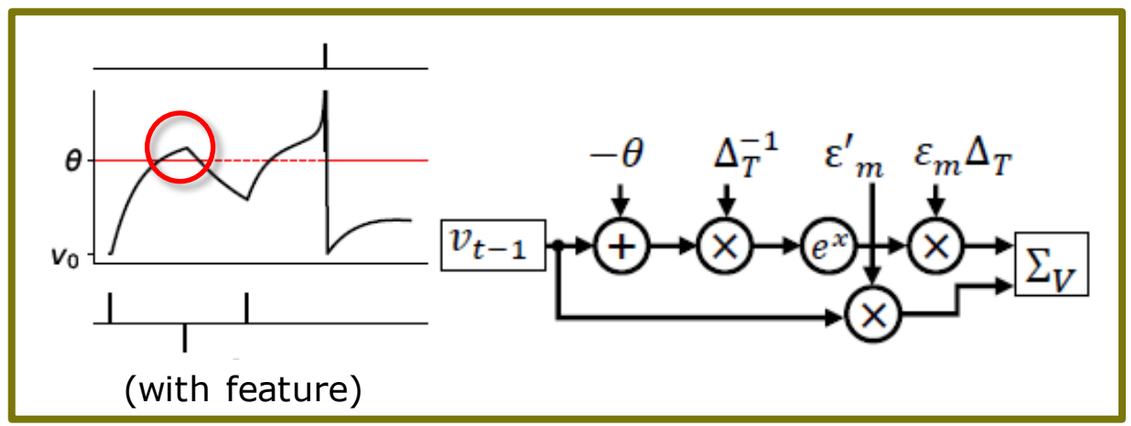

Exponential

Neuron Feature #3: **Spike-triggered Current**

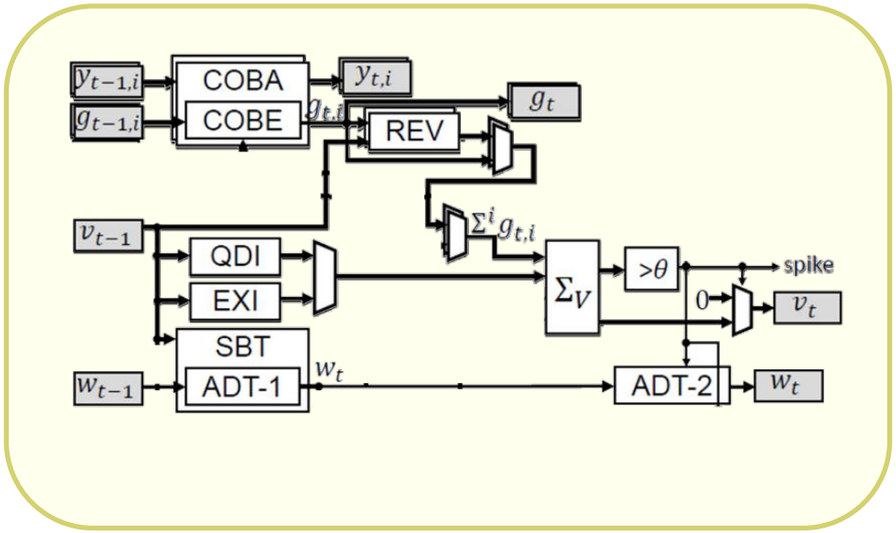

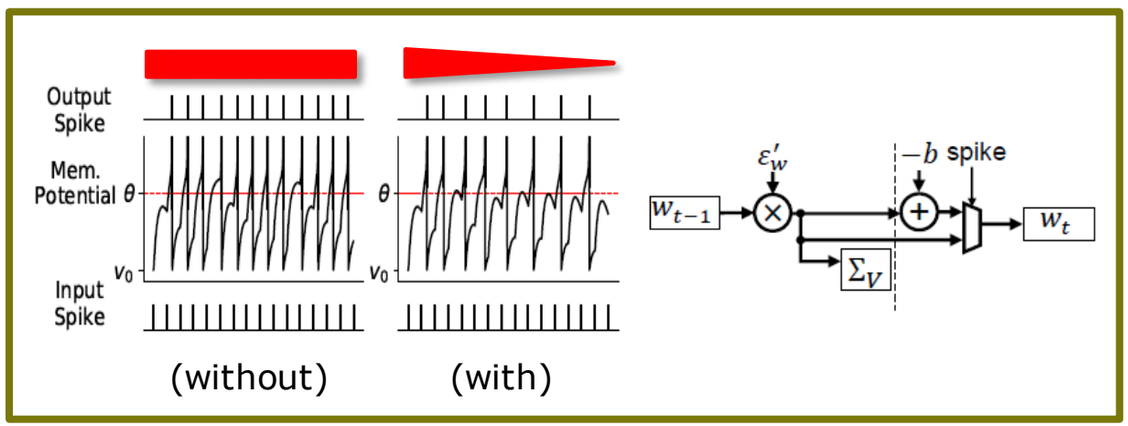

Adaptation

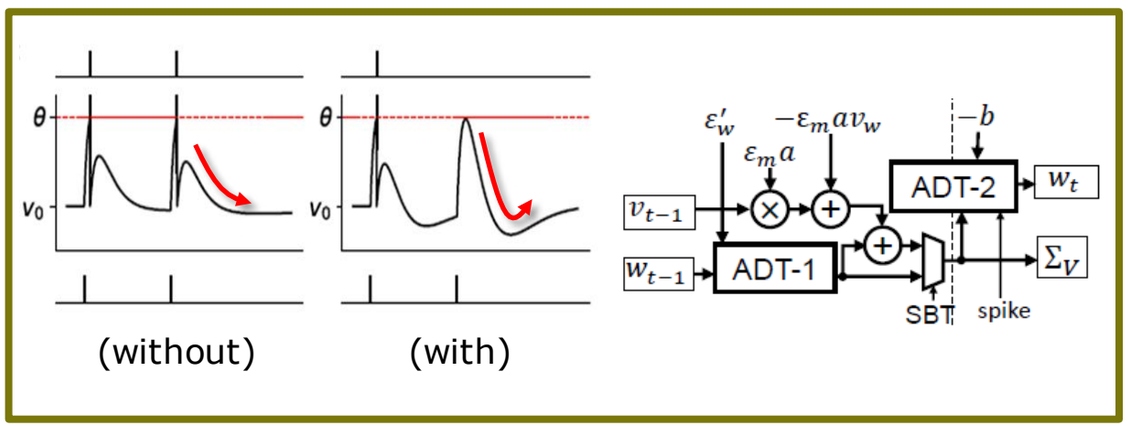

Sub-threshold Oscillation

Neuron Feature #4: **Refractory Period**

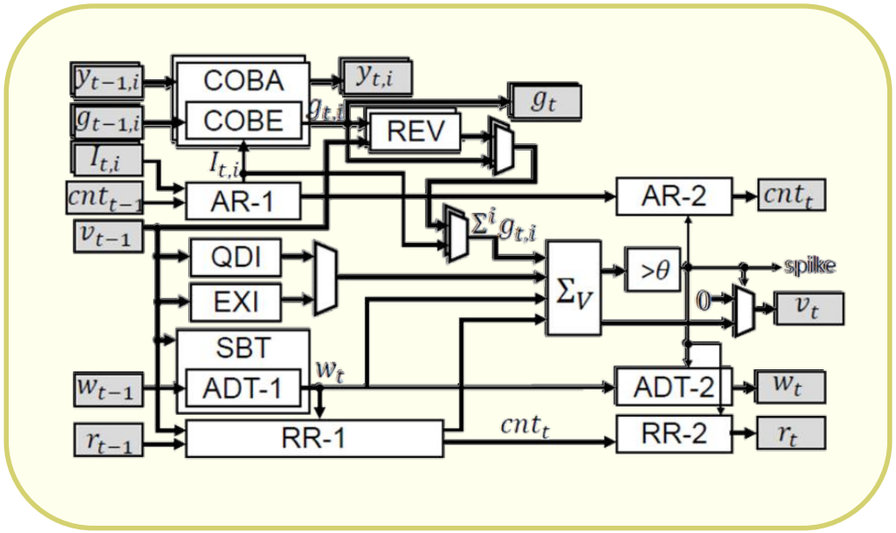

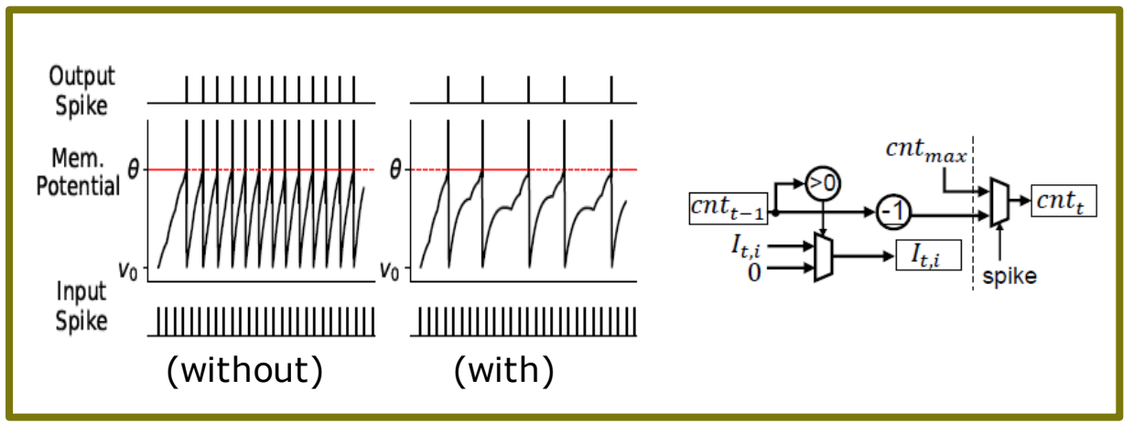

Absolute

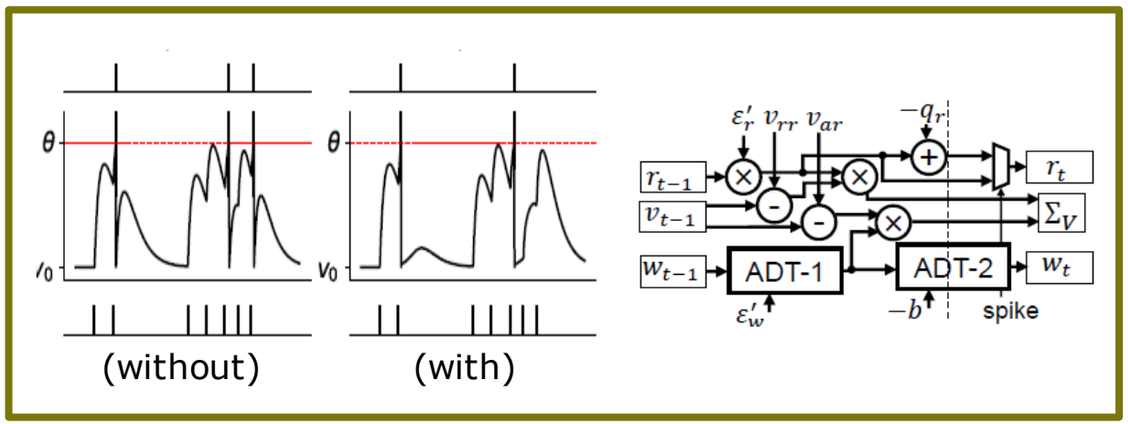

Relative

Neuron Feature: **Flexible Feature Support**

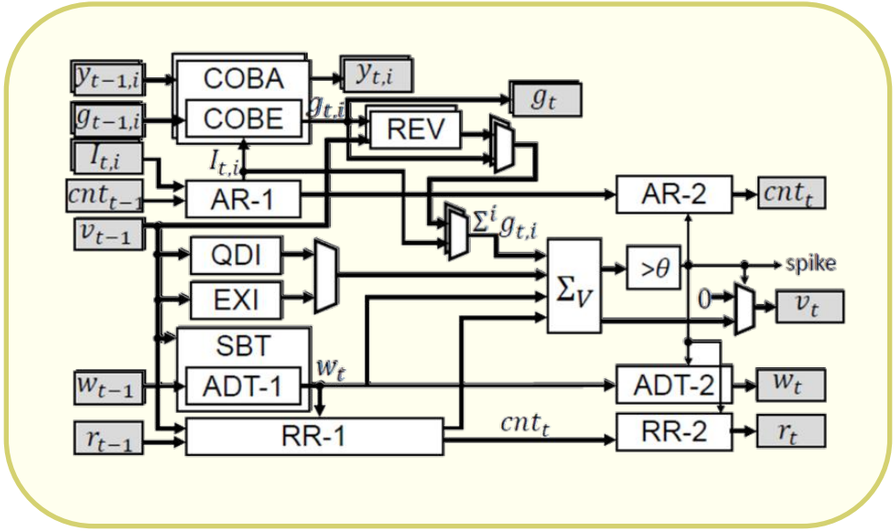

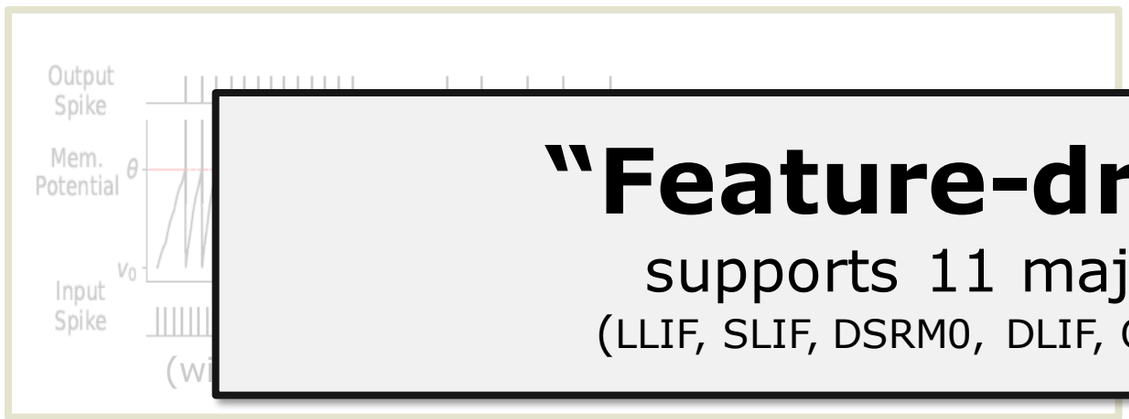

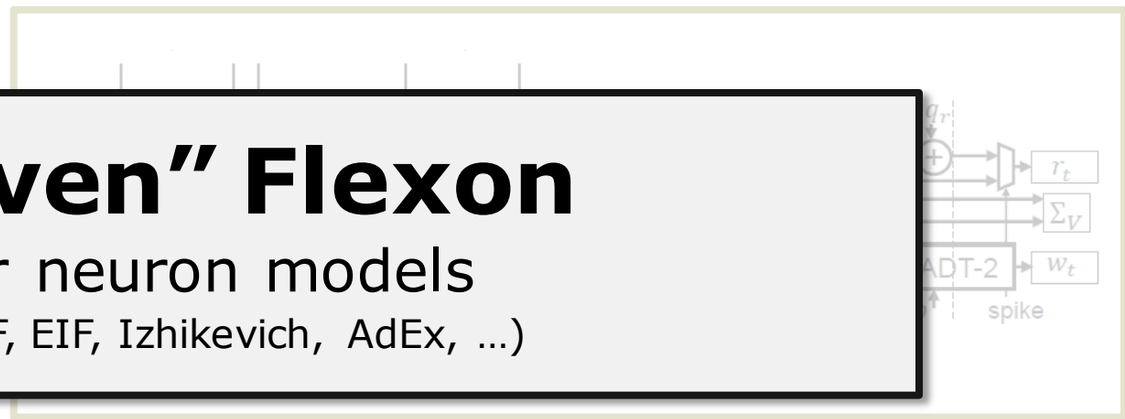

"Feature-driven" Flexon

supports 11 major neuron models

(LLIF, SLIF, DSRM0, DLIF, QIF, EIF, Izhikevich, AdEx, …)

Absolute

Relative

Evaluation (*12x Feature-driven Design*)

8 CPU + 2 GPU Representative Benchmarks

Flexon: TSMC 45nm, Synopsys Design Compiler (neuron), CACTI 6.5 (SRAM)

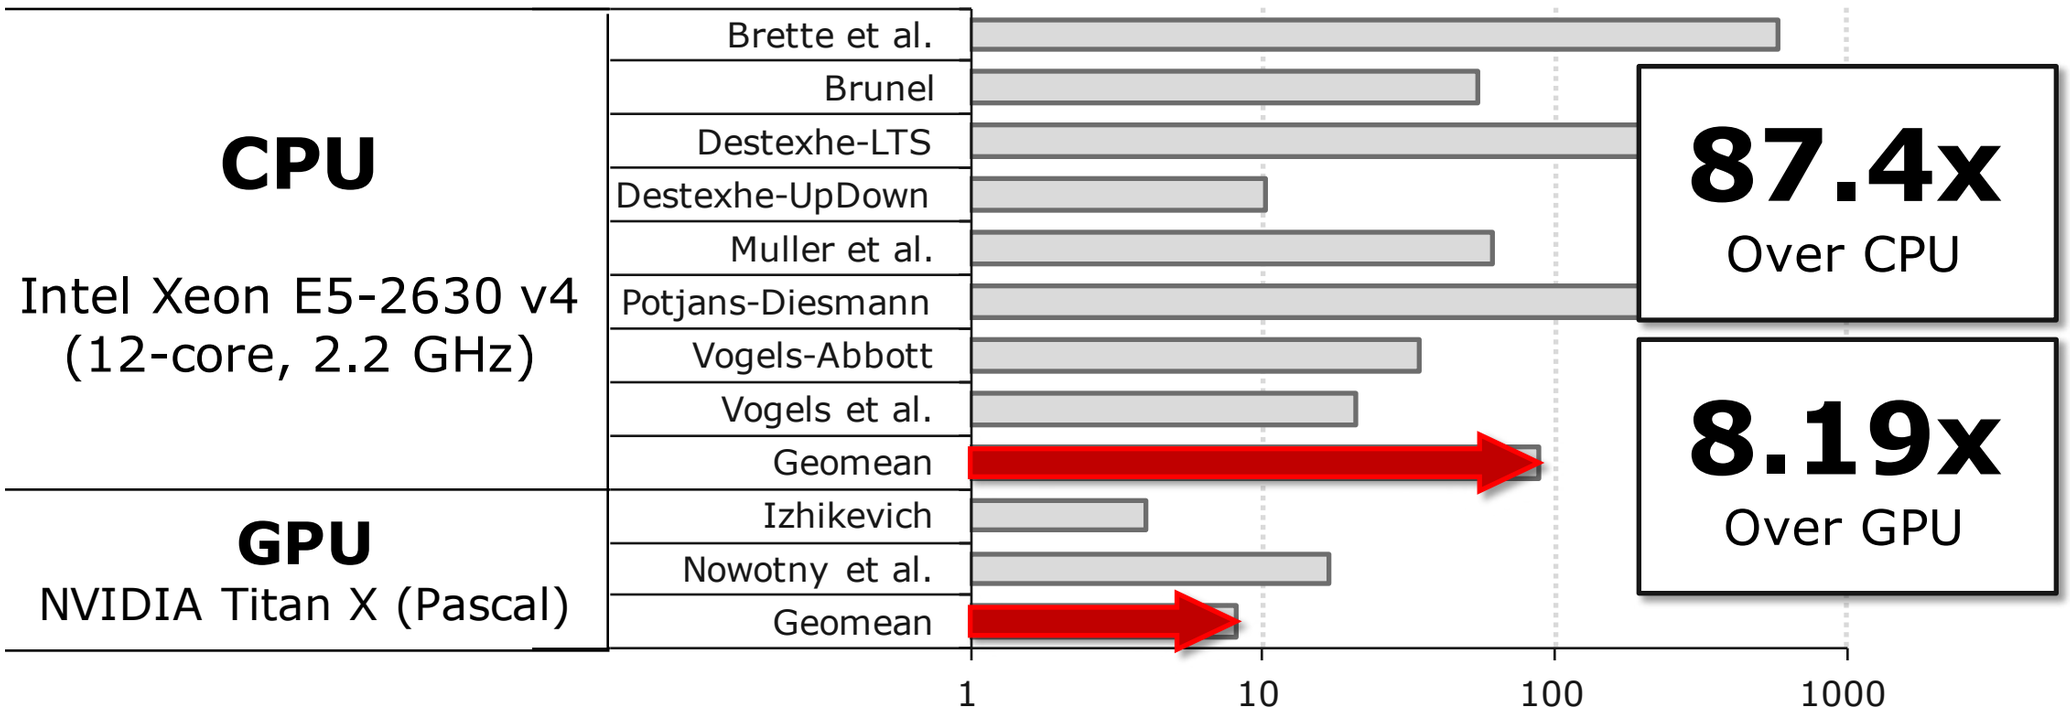

Evaluation (*12x Feature-driven Design*)

8 CPU + 2 GPU Representative Benchmarks

Flexon: TSMC 45nm, Synopsys Design Compiler (neuron), CACTI 6.5 (SRAM)

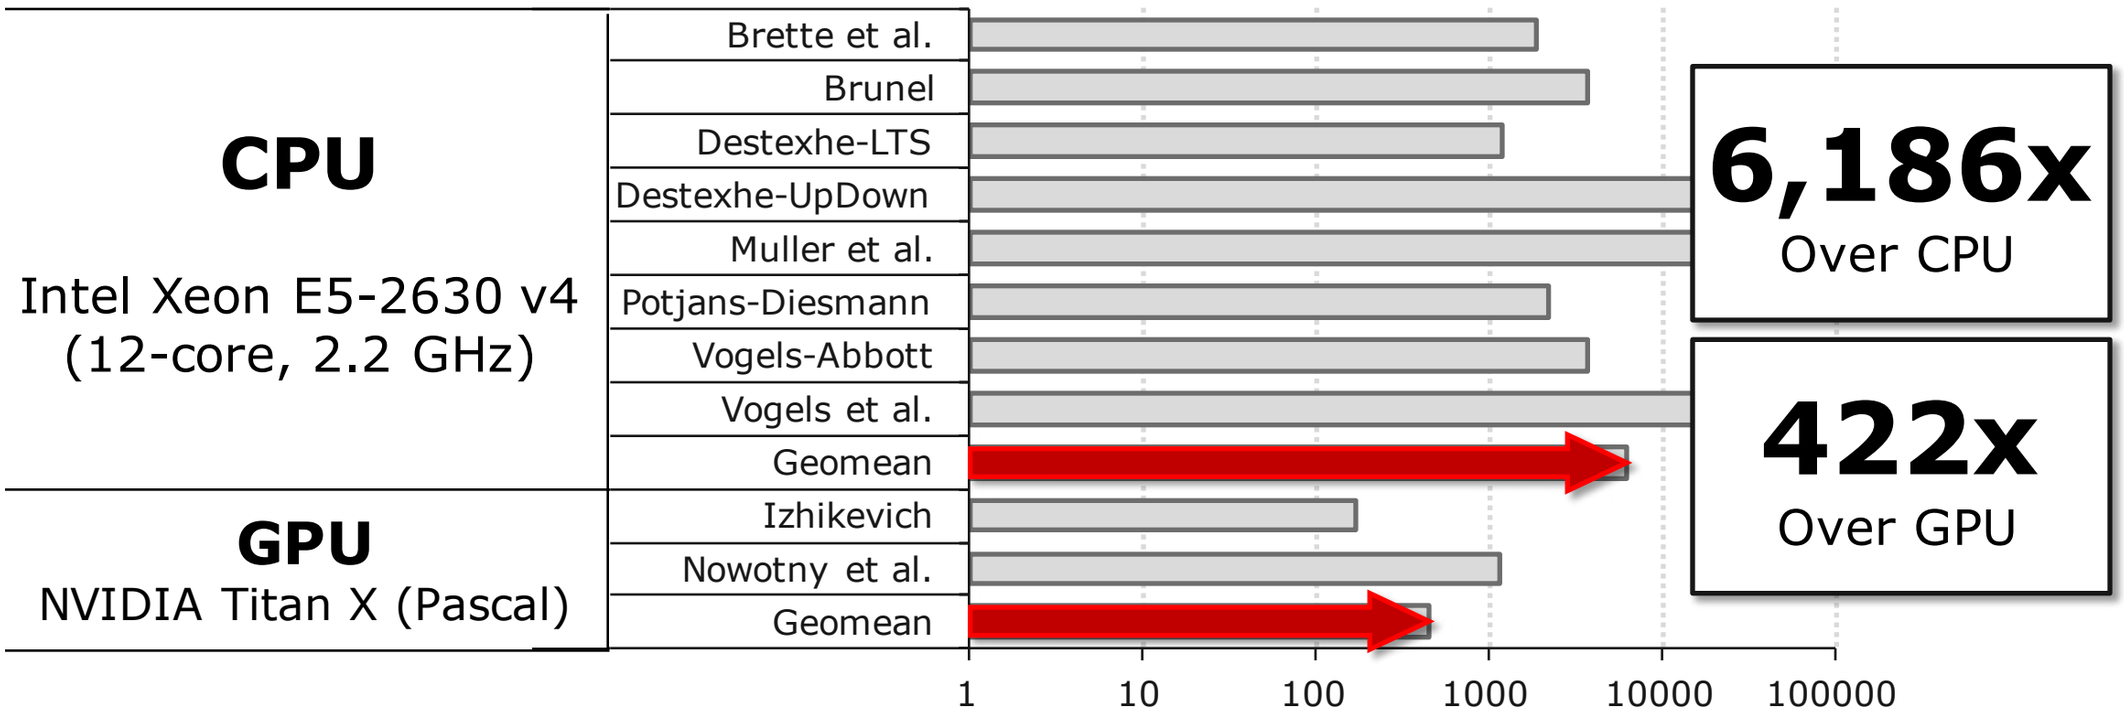

Intrinsic Space-inefficiency

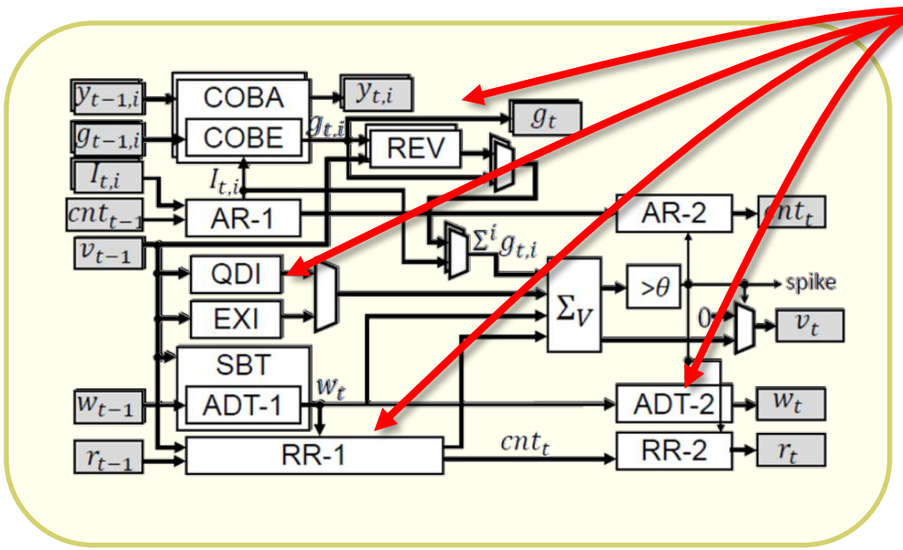

Lots of **redundant** units
(multiplier, adder, …)

“Feature-driven” Flexon

supports 11 major neuron models
(LLIF, SLIF, DSRM0, DLIF, QIF, EIF, Izhikevich, AdEx, …)

Absolute

Relative

Constructing *Spatially-folded* Flexon

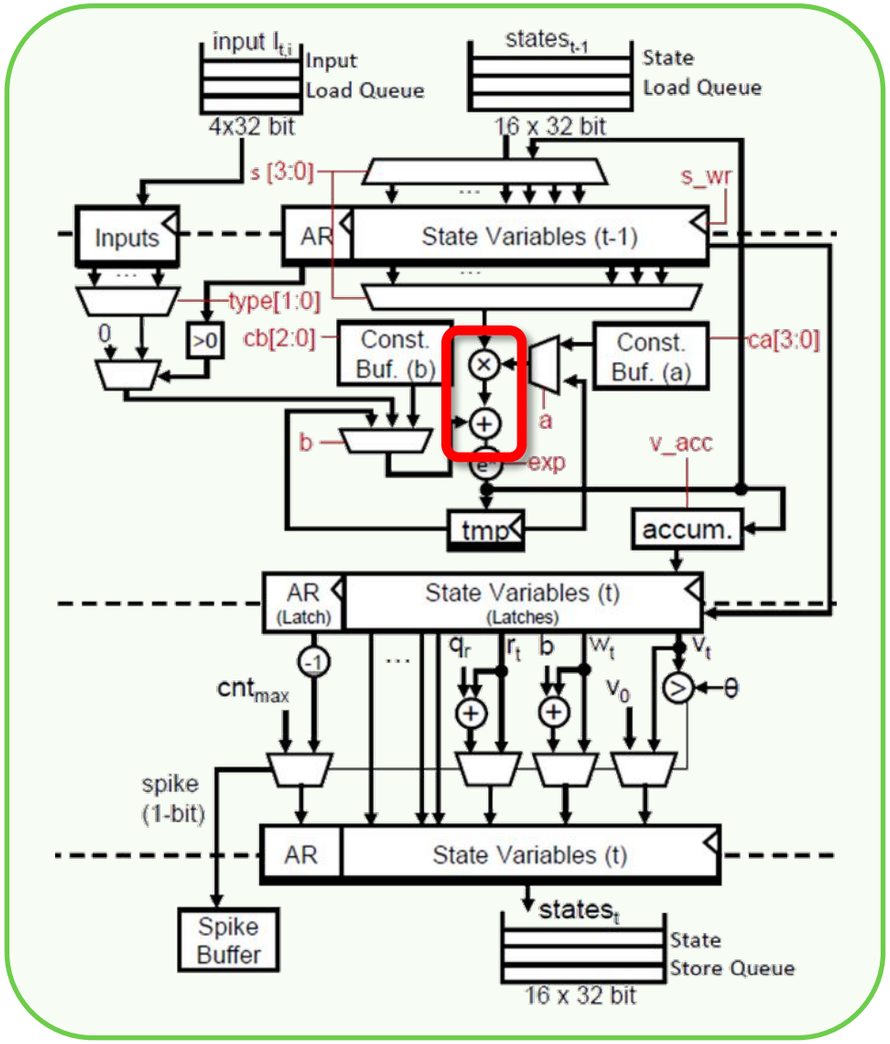

Spatially-folded design ➔ reduce area

- Remove redundant MAC operators

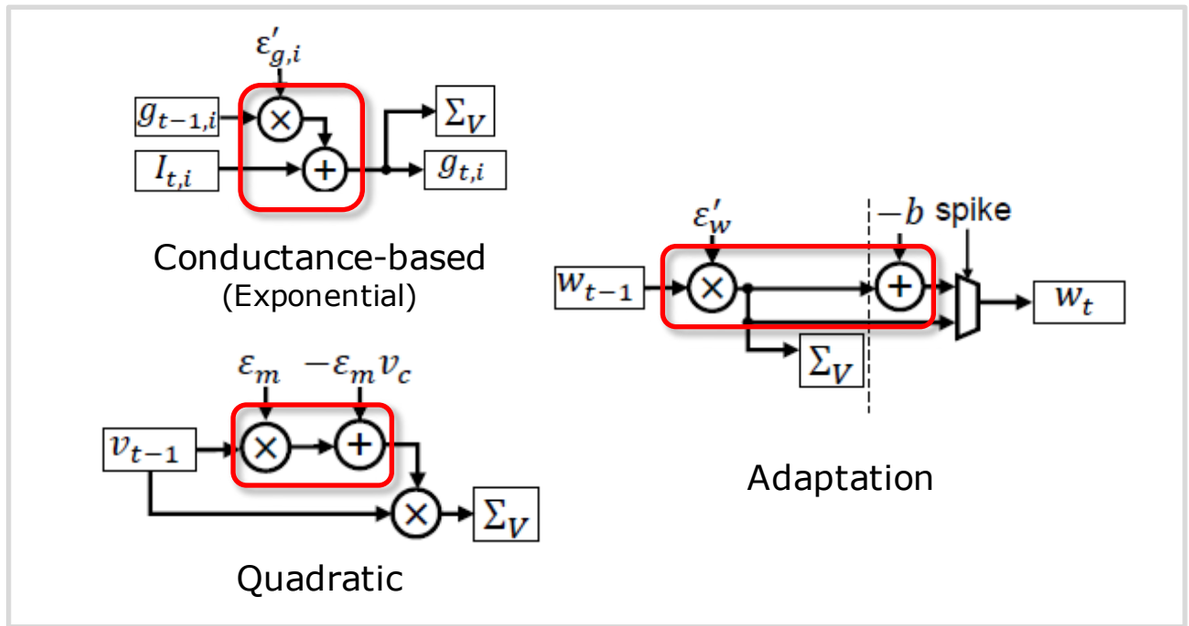

Constructing *Spatially-folded* Flexon

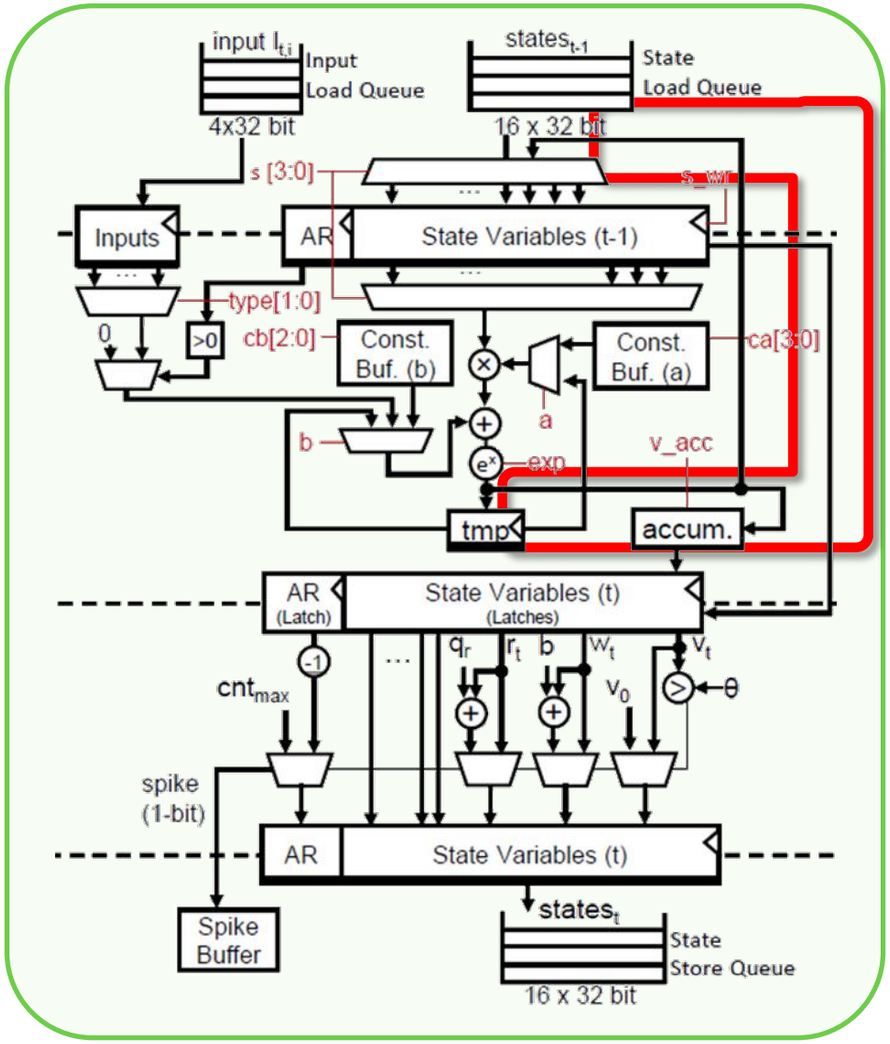

Spatially-folded design ➔ reduce area

- Remove redundant MAC operators

Modifications from the baseline

- 2-stage pipeline, multi-cycle implementation

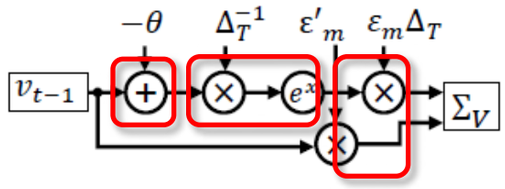

Exponential

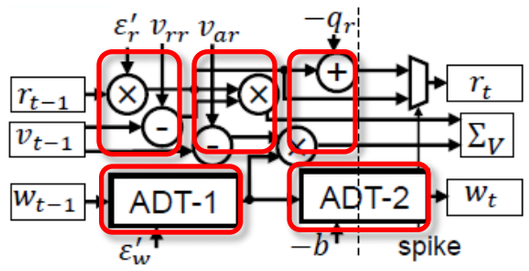

Relative
(ADT: Adaptation)

Constructing *Spatially-folded* Flexon

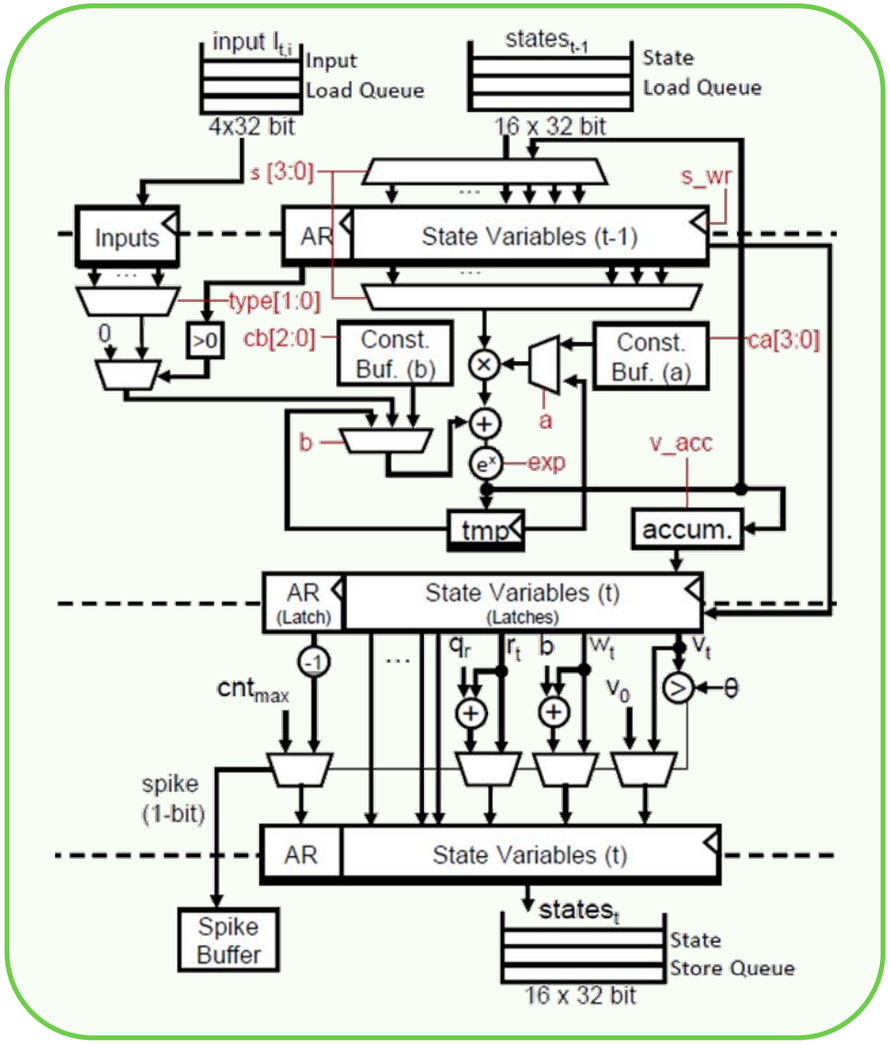

Spatially-folded design ➔ reduce area

- Remove redundant MAC operators

What we should change

- 2-stage pipeline, multi-cycle implementation

"Spatially-folded" Flexon

supports various major neuron models

6x area saving

Evaluation (*72x Spatially-folded* Design)

8 CPU + 2 GPU Representative Benchmarks

Flexon: TSMC 45nm, Synopsys Design Compiler (neuron), CACTI 6.5 (SRAM)

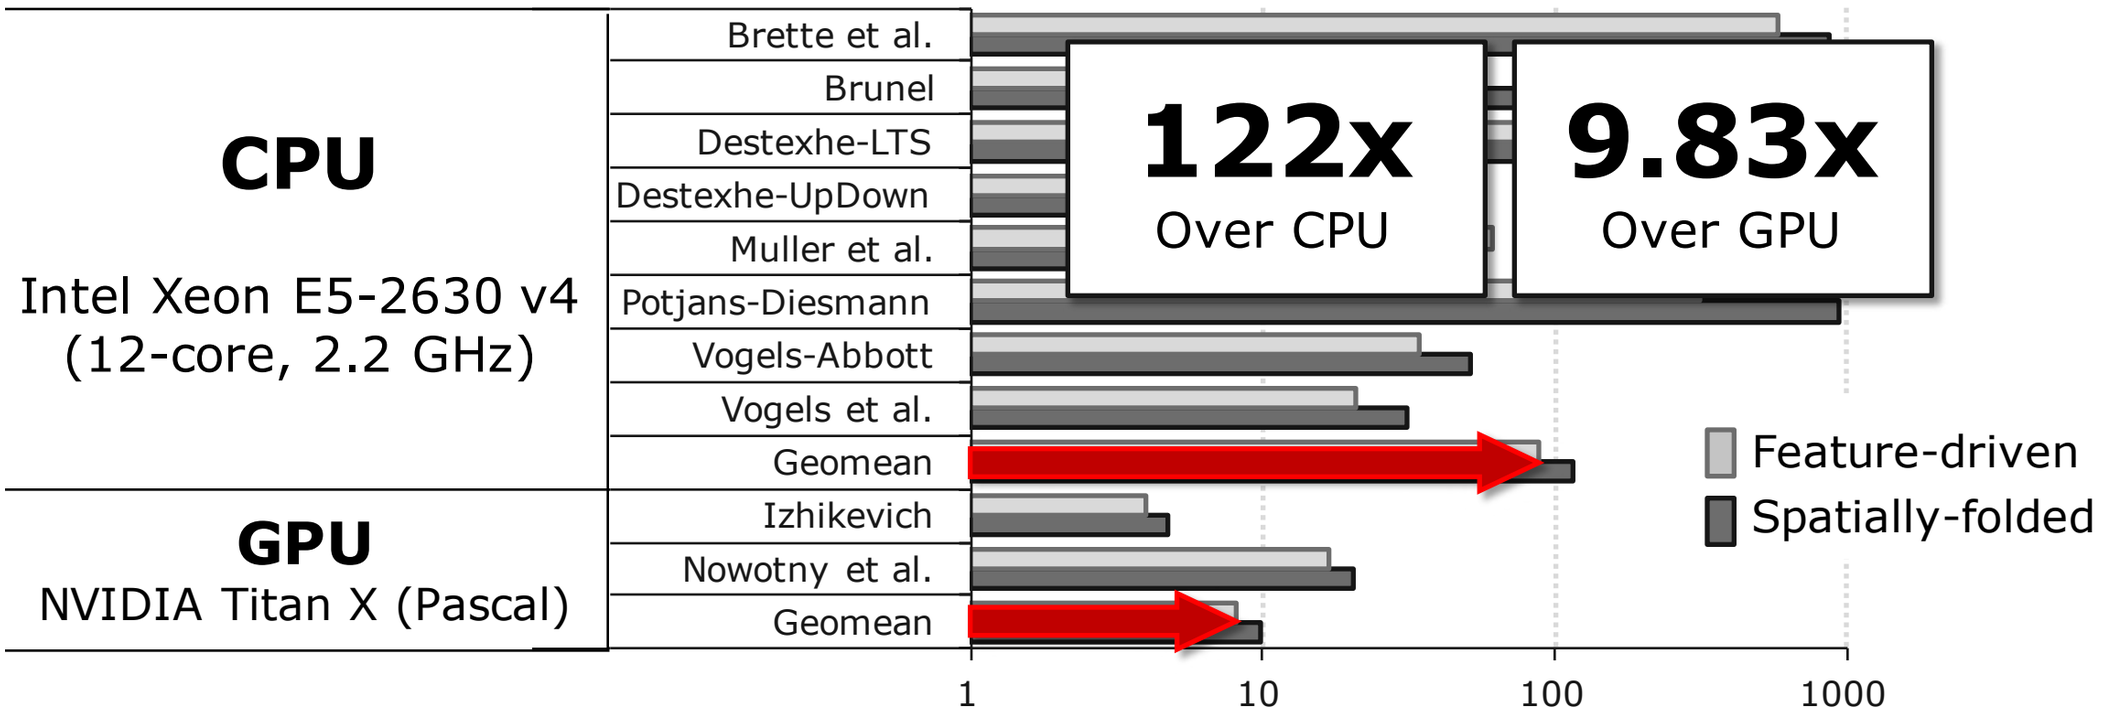

Evaluation (*72x Spatially-folded* Design)

8 CPU + 2 GPU Representative Benchmarks

Flexon: TSMC 45nm, Synopsys Design Compiler (neuron), CACTI 6.5 (SRAM)

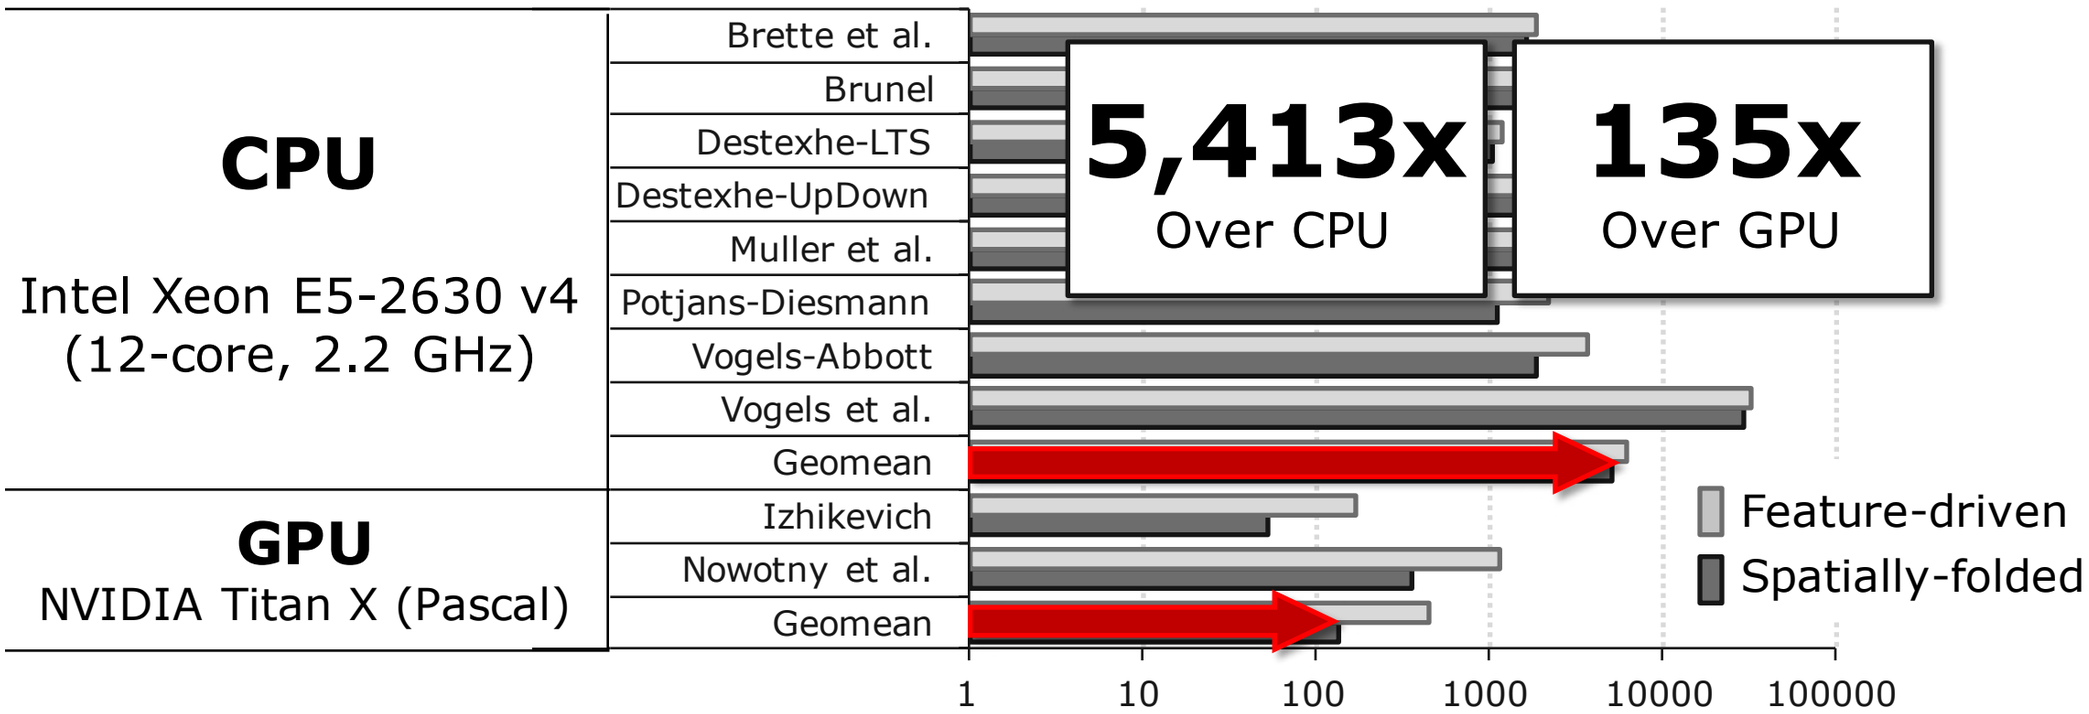

Energy Efficiency

(Normalized to the baseline)

Baseline Flexon vs. Spatially-folded Flexon

- **Baseline "Feature-driven" Flexon**
  - Fast: 87.4x over CPUs, 8.19x over GPUs
  - Energy-efficient: **6,186x** over CPUs, **422x** over GPUs

- **"Spatially-folded" Flexon**
  - Fast: **122x** over CPUs, **9.83x** over GPUs
  - Energy-efficient: 5,413x over CPUs, 135x over GPUs

Conclusion

- **Flexon** is a flexible feature-driven digital neuron design, capable of realizing various major neuron models.
  - Flexible & **power-efficient** (6,186x over CPU)

- **Spatially-folded Flexon** makes features share units, reducing 6x circuit area.
  - Flexible & **fast** when integrated (122x over CPU)

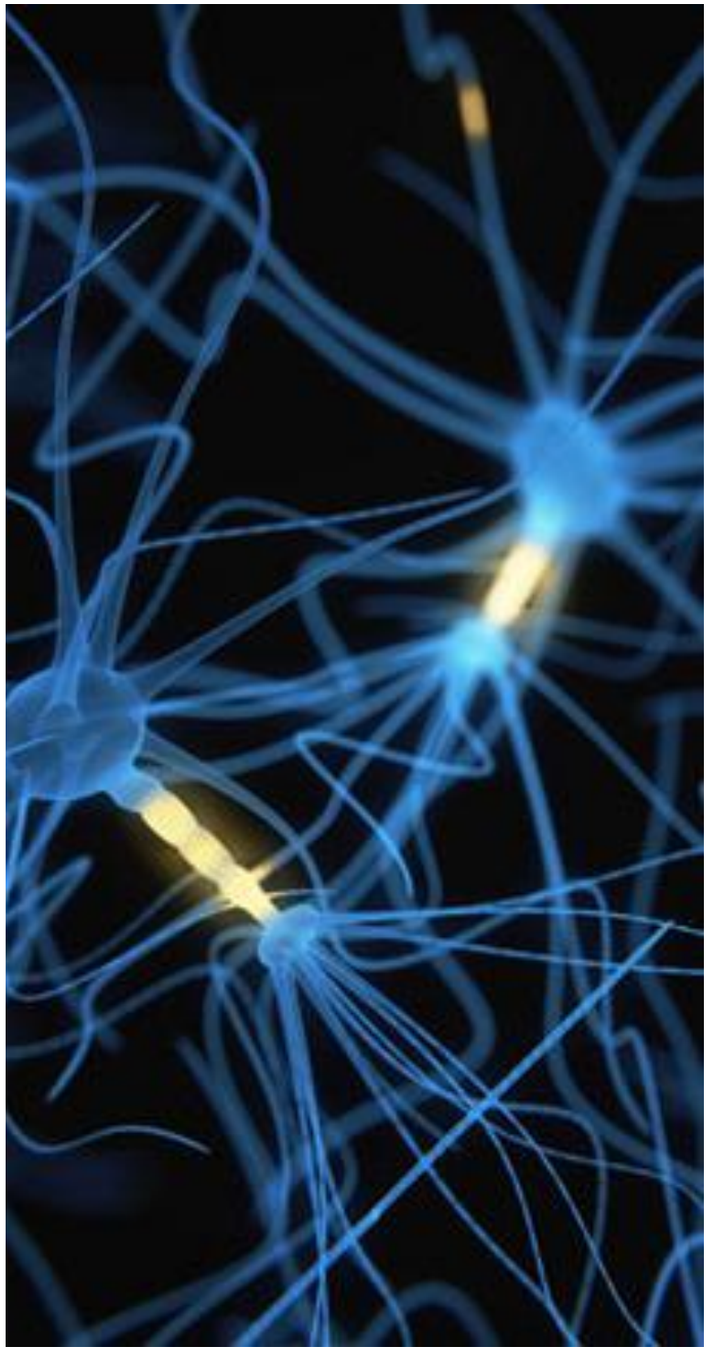

HPCS
High Performance Computer System Lab

Flexon

A Flexible Digital Neuron for Efficient
Spiking Neural Network Simulations

Thank you for listening