




SIMBA: SCALING DEEP-LEARNING INFERENCE WITH MULTI-CHIP-MODULE-BASED ARCHITECTURE


Yakun Sophia Shao, Jason Clemons, Rangharajan Venkatesan, Brian Zimmer, Matthew Fojtik, Nan Jiang, Ben Keller, Alicia Klinefelter, Nathaniel Pinckney, Priyanka Raina, Stephen G. Tell, Yanqing Zhang, William J. Dally, Joel S. Emer, C. Thomas Gray, Brucek Khailany & Stephen W. Keckler

VAST WORLD OF AI INFERENCE

Creating A Massive Market Opportunity



MOBILE DEVICES



EMBEDDED COMPUTERS



DATACENTER COMPUTERS

SCALABLE INFERENCE ACCELERATORS

Challenges and Opportunities

Motivation

- Need for fast and efficient inference accelerators from mobile to datacenter.

Challenge

- High design cost of building unique hardware for each design target.

Opportunities

- Deep learning inference is intrinsically scalable with abundant parallelism.
- Recent advances in package-level integration for multi-chip-module-based designs.

THE MULTI-CHIP-MODULE APPROACH


Challenges and Opportunities

Advantages:

- Build systems larger than reticle limit
- Smaller chips are cheaper to design
- Smaller chips have higher yield
- Faster time-to-market

Challenges:

- Area, energy, and latency for chip-to-chip communication



Ref: Zimmer et al., VLSI 2019

SIMBA: SCALING INFERENCE WITH MULTI-CHIP-MODULE-BASED ARCHITECTURE


Simba Testchip:

- Package and chiplet architecture
- Processing element design
- Baseline uniform tiling across chiplets and PEs




Simba Characterization:

- Comparison with GPUs
- NoP bandwidth sensitivity
- NoP latency sensitivity



Simba NoP-Aware Tiling:


- Non-uniform work partitioning
- Communication-aware data placement
- Cross-layer pipelining



SIMBA: SCALING INFERENCE WITH MULTI-CHIP-MODULE-BASED ARCHITECTURE


Simba Testchip:

- Package and chiplet architecture
- Processing element design
- Baseline uniform tiling across chiplets and PEs




Simba Characterization:

- Comparison with GPUs
- NoP bandwidth sensitivity
- NoP latency sensitivity



Simba NoP-Aware Tiling:

- Non-uniform work partitioning
- Communication-aware data placement
- Cross-layer pipelining



SIMBA: SCALABLE MCM-BASED ARCHITECTURE

Simba Package and Chiplet

Package and chiplet spec

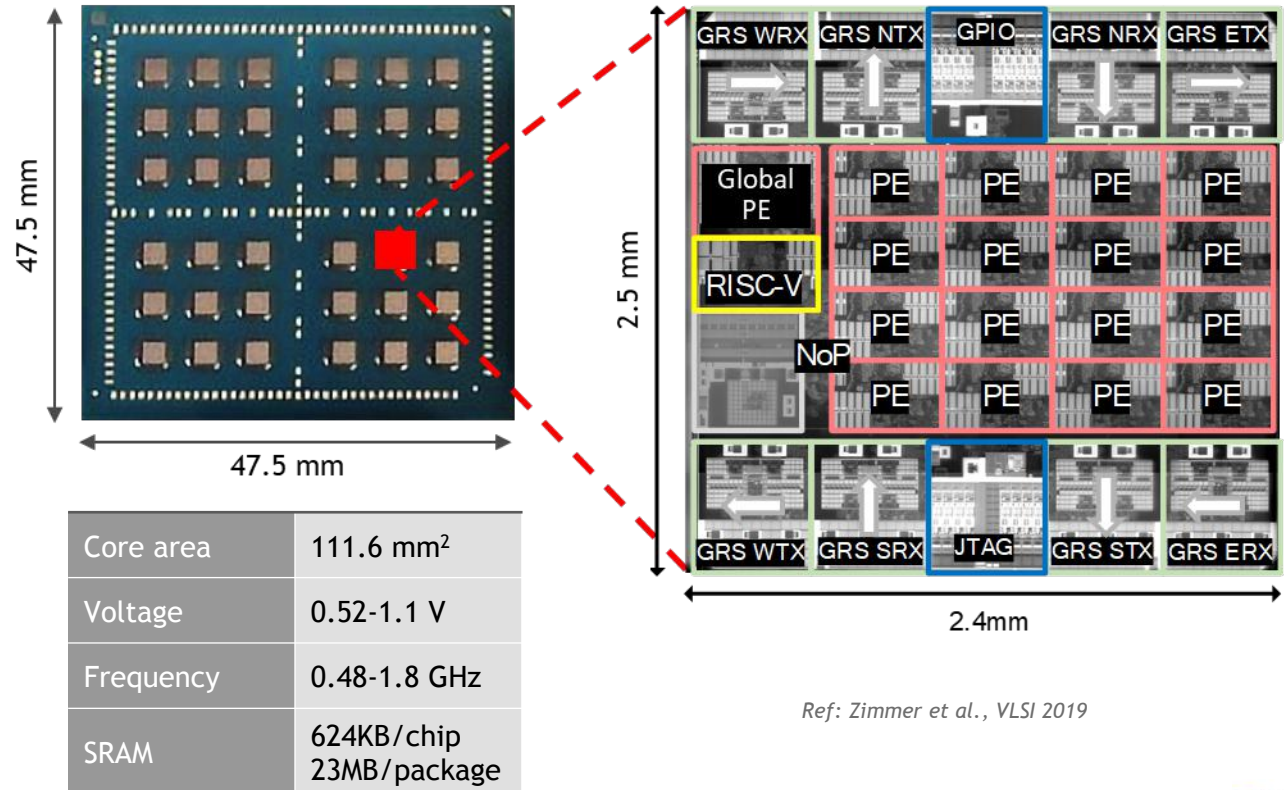
6mm² chiplet in TSMC 16nm
36 chiplets/package

Chip-to-chip interconnect

Ground-Referenced Signaling

Efficient compute tiles


128 TOPS
0.11 pJ/Op
8-bit integer datapath



Ref: Zimmer et al., VLSI 2019

SIMBA: NON-UNIFORM COMMUNICATION

Network-on-Chip (NoC) and Network-on-Package (NoP)



NETWORK-ON-CHIP (NoC)

4x5 mesh topology connects 16 PEs, one Global PE, and one RISC-V.

Cut-through routing with multicast support.

10ns per hop, 68GB/s/PE

NETWORK-ON-PACKAGE (NoP)

6x6 mesh topology connects 36 chiplets in package.

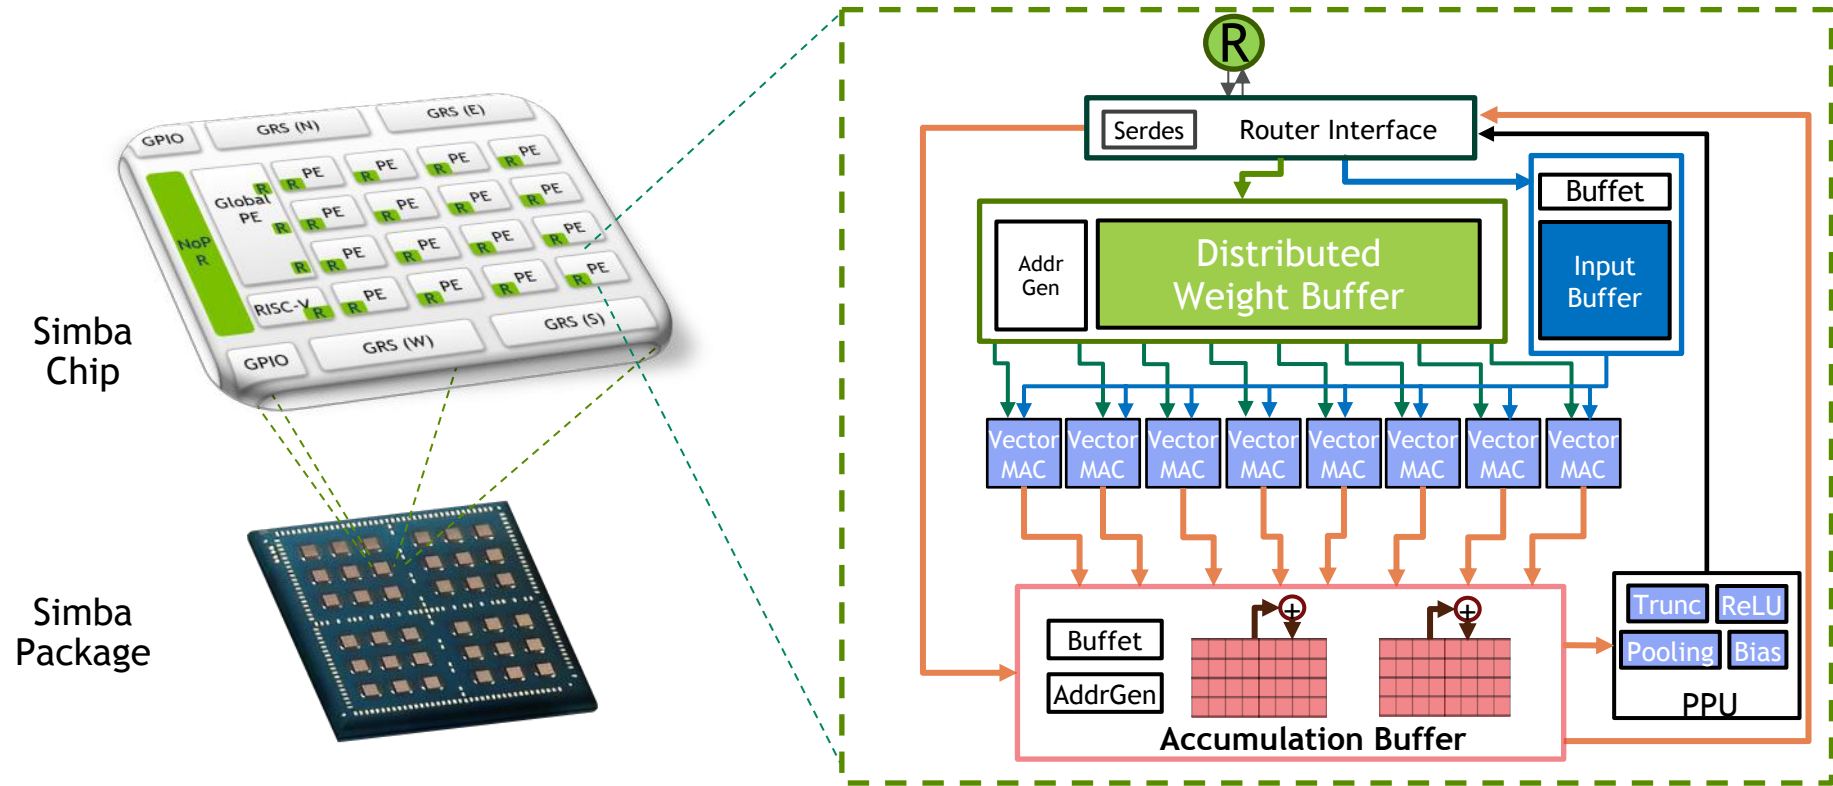
A NoP router per chiplet.

Configurable routing to avoid bad links/chiplets.

20ns per hop, 100GB/s/chiplet


SIMBA: SCALABLE MCM-BASED ARCHITECTURE

Spatial Architecture with Distributed Memory



BASELINE: UNIFORM TILING ACROSS NOC/NOP


Exploiting Model and Data Parallelism



SIMBA: SCALING INFERENCE WITH MULTI-CHIP-MODULE-BASED ARCHITECTURE


Simba Testchip:

- Package and chiplet architecture
- Processing element design
- Baseline uniform tiling across chiplets and PEs




Simba Characterization:

- Comparison with GPUs
- NoP bandwidth sensitivity
- NoP latency sensitivity



Simba NoP-Aware Tiling:

- Non-uniform work partitioning
- Communication-aware data placement
- Cross-layer pipelining



SIMBA MEASUREMENT SETUP

Measurements begin after weights and activations are loaded from FPGA DRAM

Weights are loaded to PE memory
Activations are loaded to Global PE

Operating points


Voltage: 0.72 V

Frequency: 1.03 GHz

GRS Bandwidth: 11 Gbps/pin


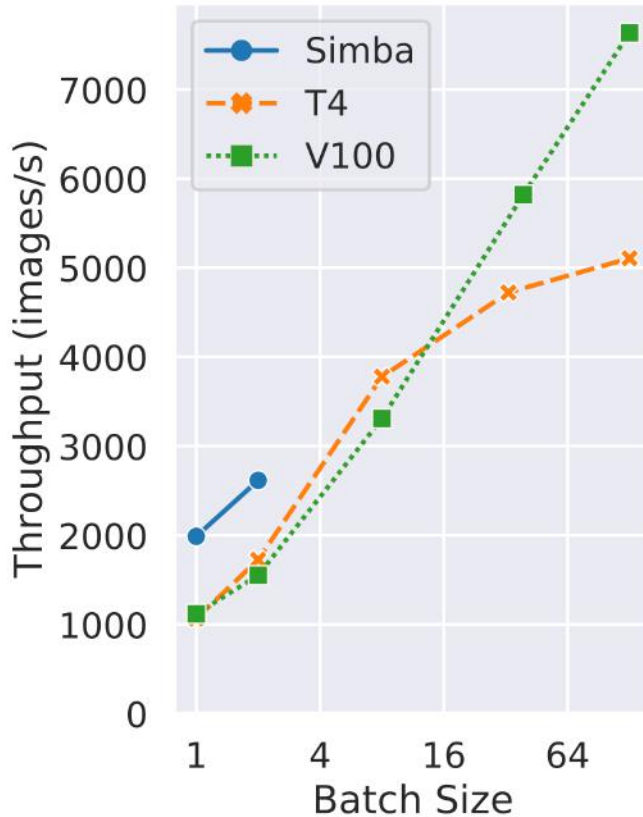
Networks

DriveNet, AlexNet, ResNet-50



SIMBA CHARACTERIZATION


Comparison with GPUs running ResNet-50



SIMBA CHARACTERIZATION

Layer Sensitivity

- Running three ResNet-50 layers across different number of chiplets.
- Increasing the number of active chiplets does not always translate to performance gains.
- The cost of communication hinders the ability to exploit parallelism.



SIMBA CHARACTERIZATION

NoP Bandwidth Sensitivity


- Running two ResNet-50 layers across 32 chiplets with different NoP bandwidths by adjusting the ratio between NoP and PE frequencies.
- End-to-end performance is sensitive to NoP bandwidth, especially for applications with a significant amount of communication.



SIMBA CHARACTERIZATION

NoP Latency Sensitivity


- Mapping res4a_branch1 to four chiplets with different data placements (shown in the X-axis)
- Inter-chiplet traffic type:
 - Input activation multicast
 - Output activation writeback
- Communication latency plays a key role in achieving good performance/efficiency in a large-scale system.



SIMBA: SCALING INFERENCE WITH MULTI-CHIP-MODULE-BASED ARCHITECTURE


Simba Testchip:

- Package and chiplet architecture
- Processing element design
- Baseline uniform tiling across chiplets and PEs




Simba Characterization:

- Comparison with GPUs
- NoP bandwidth sensitivity
- NoP latency sensitivity



Simba NoP-Aware Tiling:

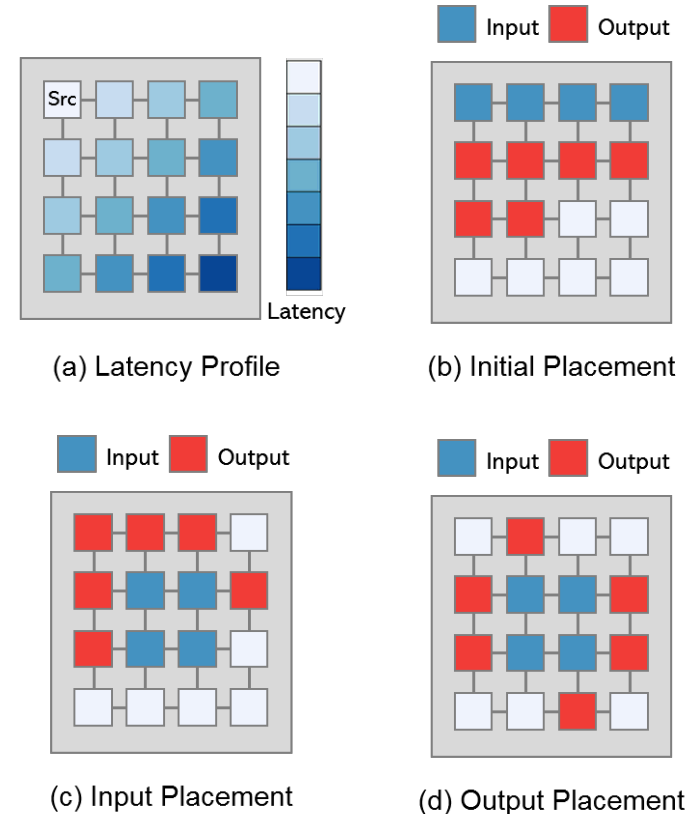
- Non-uniform work partitioning
- Communication-aware data placement
- Cross-layer pipelining



SIMBA NOP-AWARE TILING

Communication-Aware Data Placement


- Communication latency is highly sensitive to the physical location of data, which is explicitly managed by programmers.
- We use an iterative algorithm to place the input and output activations to minimize the number of NoP hops.
- We keep the same tiling in the updated placements to reduce the number of variables.



SIMBA NOP-AWARE TILING

Communication-Aware Data Placement


- Communication latency is highly sensitive to the physical location of data, which is explicitly managed by programmers.
- We use an iterative algorithm to place the input and output activations to minimize the number of NoP hops.
- We keep the same tiling in the updated placements to reduce the number of variables.



SIMBA: SCALING INFERENCE WITH MULTI-CHIP-MODULE-BASED ARCHITECTURE


Simba Testchip:

- Package and chiplet architecture
- Processing element design
- Baseline uniform tiling across chiplets and PEs




Simba Characterization:


- Comparison with GPUs
- NoP bandwidth sensitivity
- NoP latency sensitivity



Simba NoP-Aware Tiling:

- Non-uniform work partitioning
- Communication-aware data placement
- Cross-layer pipelining





**SIMBA: SCALING DEEP-LEARNING INFERENCE WITH
MULTI-CHIP-MODULE-BASED ARCHITECTURE**



nvidia®